

Lecture 6

Relational Algebra

COMP3278A

Introduction to Database Management Systems

Dr. Ping Luo

Email : pluo@cs.hku.hk



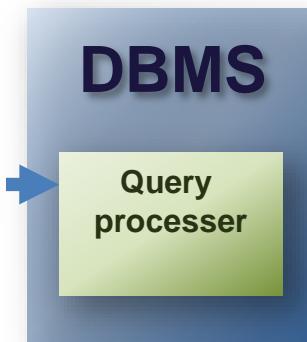
Department of Computer Science, The University of Hong Kong

Acknowledgement: Dr Chui Chun Kit

Motivation

- Find the dept. names where employees named Smith work.

```
SELECT D.name  
FROM Employees E, Works_in W, Departments D  
WHERE E.name = 'Smith' AND  
E.employee_id = W.employee_id AND  
W.department_id = D.department_id;
```



How does the DBMS execute this SQL query?

e.g. Join which two tables first?

Which constraint is applied first?



Employees

employee_id	name	salary
1	Jones	26000
2	Smith	28000
3	Parker	35000
4	Smith	24000

Works_in

employee_id	department_id	since
1	1	2001-1-1
2	1	2002-4-1
2	2	2005-2-2
3	3	2003-1-1
4	3	2005-1-1

Departments

department_id	name	budget
1	Toys	122000
2	Tools	239000
3	Food	100000

Relational Algebra

- Relational Algebra is similar to normal algebra (as in $2+3*x-y$), except it uses **relations (tables)** as **operand**, and a new set of **operators**.
- The inner, lower-level operations of a relational DBMS are (or are similar to) relational algebra operations.
- We need to know about relational algebra to understand query execution and optimization in a relational DBMS.

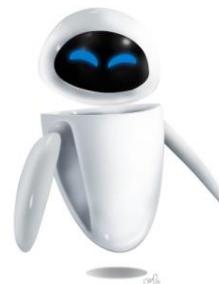
Section 1

Basic operators

Basic operators

- Select (σ)
- Project (π)
- Union (\cup)
- Set difference (-)
- Cartesian product (\times)
- Rename (ρ)

That is to say, there should be programs (or functions) implemented in the DBMS for each of these relational operators



Selection

- $\sigma_p(R) = \{ t \mid t \in R \wedge p(t) \}$

- Example

- Consider the relation **Author** (authorID, name, date of birth)
- **Select** all authors called “May”.

```
SELECT *
FROM Author
WHERE name = "May";
```

SQL

Query
processor

$\sigma_{\text{name}=\text{"May"}}(\text{Author})$

Relational algebra

Author

authorID	name	date of birth
101	May	Nov 16
102	Bonnie	Jan 15
103	May	Jul 11
104	Raymond	Apr 30
105	Tiffany	Oct 10

$\sigma_{\text{name}=\text{"May"}}(\text{Author})$

authorID	name	date of birth
101	May	Nov 16
103	May	Jul 11

Projection



$\pi_{A_1, A_2, \dots, A_k}(R)$

The result = relation over the k attributes A_1, A_2, \dots, A_k obtained from R by erasing the columns that are not listed and eliminating duplicate rows

- A copy of R with only listed attributes A_1 to A_k .



Example

- Consider the relation **Book** (bookID, title, publisher)
- Report only the **bookID** and **title** of all the books.

```
SELECT bookID, title  
FROM Book;
```

SQL

Query
processor

$\pi_{\text{bookID}, \text{title}}(\text{Book})$

Relational algebra

Book

bookID	title	publisher
115	Stuffy doll	ABC
116	Angel's feather	MTG
117	Little girl	MGH
118	Myr	ABC

$\pi_{\text{bookID}, \text{title}}(\text{Book})$

bookID	title
115	Stuffy doll
116	Angel's feather
117	Little girl
118	Myr

Union

- $R \cup S = \{ t \mid t \in R \vee t \in S \}$
- R and S must have the **same number of attributes** and **attribute data types are compatible**.
- Example
 - Find the name of all products in Audio_CD and DVD tables.

```
SELECT name  
FROM Audio_CD  
UNION  
SELECT name  
FROM DVD;
```

SQL



$\pi_{name}(\text{Audio_CD}) \cup \pi_{name}(\text{DVD})$

Relational algebra

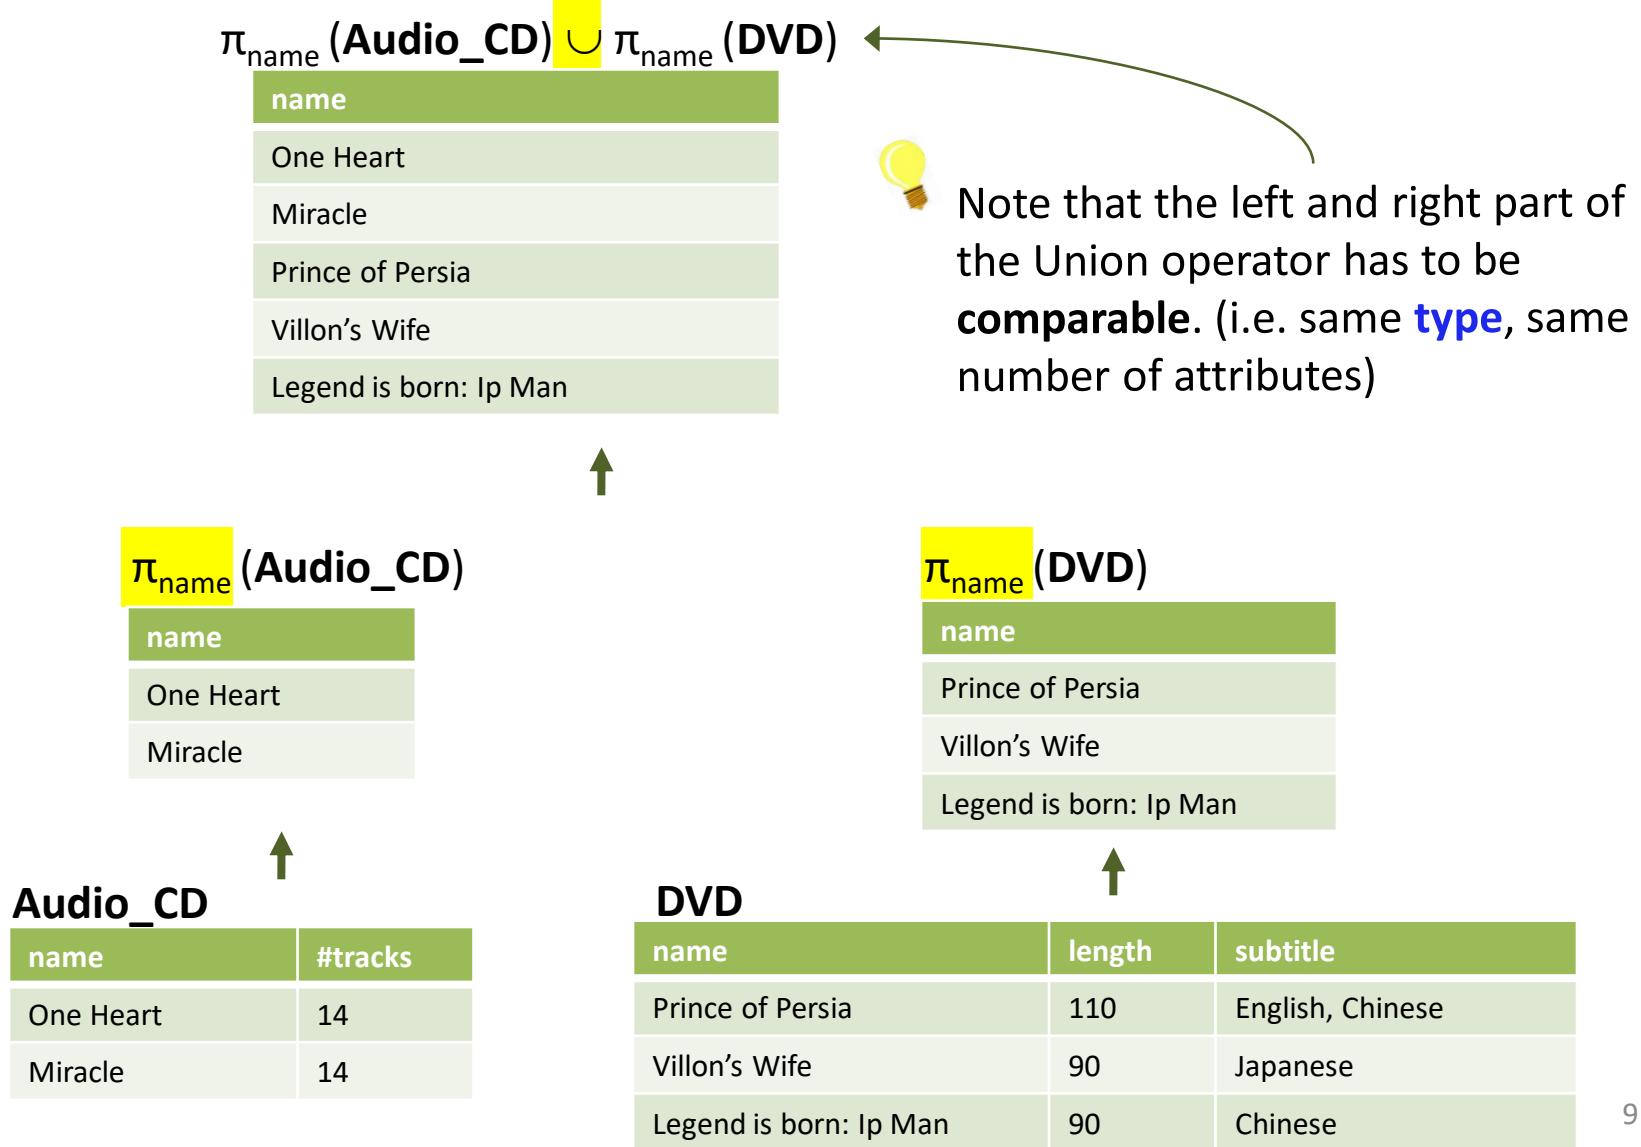
Audio_CD

name	#tracks
One Heart	14
Miracle	14

DVD

name	length	subtitle
Prince of Persia	110	English, Chinese
Villon's Wife	90	Japanese
Legend is born: Ip Man	90	Chinese

Union



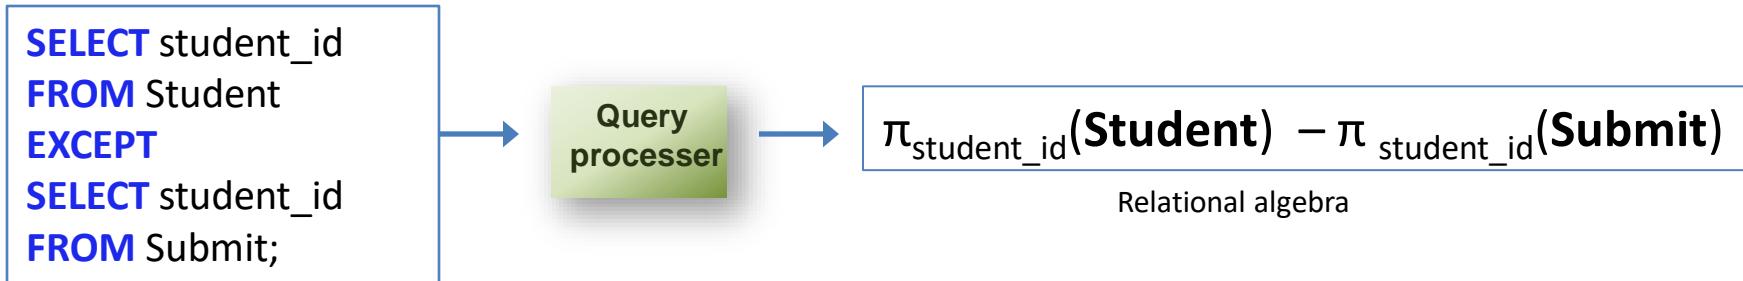
Set difference

● $R - S = \{ t \mid t \in R \wedge t \notin S \}$

- R and S must have the **same number of attributes** and **attribute data types are compatible**.

Example

- Find the ID of the students who haven't submitted the assignment.



Student			
student_id	name	gender	major
123	Kit	M	CS
456	Yvonne	F	CS
789	Paul	M	CS

Submit		
student_id	assignment_id	date
456	1	28/9
789	1	25/9

Set difference

$\pi_{\text{student_id}}(\text{Student}) - \pi_{\text{student_id}}(\text{Submit})$

student_id
123



Note that the left and right part of the Set difference operator has to be **comparable**. (i.e. same **type**, same number of attributes)

$\pi_{\text{student_id}}(\text{Student})$

student_id
123
456
789

$\pi_{\text{student_id}}(\text{Submit})$

student_id
456
789

Student

student_id	name	gender	major
123	Kit	M	CS
456	Yvonne	F	CS
789	Paul	M	CS

Submit

student_id	assignment_id	date
456	1	28/9
789	1	25/9

Cartesian product

- $R \times S = \{ t q \mid t \in R \wedge q \in S \}$
 - No attributes with a common name in R and S.
- Example
- Display the date of the tutorials of the course “Introduction to Database Management Systems”.

```
SELECT Tutorial.date  
FROM Course, Tutorial  
WHERE  
Coruse.name="Introduction to  
Database Management Systems" AND  
Course.course_id = Tutorial.course_id;
```

SQL

Query processor

$$\pi_{\text{Tutorial.date}} (\sigma_{\text{Course.name}=\text{"Introduction to Database Management Systems"} (\sigma_{\text{Course.course_id}=\text{Tutorial.course_id}} (\text{Course} \times \text{Tutorial}))})$$

Relational algebra

Course

course_id	name
c1119	Data Structures and Algorithms
c0278a	Introduction to Database Management Systems

Tutorial

tutorial_id	course_id	date
1	c1119	5/9
1	c0278a	7/9
2	c0278a	15/9

Cartesian product

Course × Tutorial

Course.course_id	Course.name	Tutorial.tutorial_id	Tutorial.course_id	Tutorial.date
c1119	Data Structures and Algorithms	1	c1119	5/9
c1119	Data Structures and Algorithms	1	c0278a	7/9
c1119	Data Structures and Algorithms	2	c0278a	15/9
c0278a	Introduction to Database Management Systems	1	c1119	5/9
c0278a	Introduction to Database Management Systems	1	c0278a	7/9
c0278a	Introduction to Database Management Systems	2	c0278a	15/9



Course

course_id	name
c1119	Data Structures and Algorithms
c0278a	Introduction to Database Management Systems

Tutorial

tutorial_id	course_id	date
1	c1119	5/9
1	c0278a	7/9
2	c0278a	15/9

Cartesian product

Course × Tutorial

Course.course_id	Course.name	Tutorial.tutorial_id	Tutorial.course_id	Tutorial.date
c1119	Data Structures and Algorithms	1	c1119	5/9
c1119	Data Structures and Algorithms	1	c0278a	7/9
c1119	Data Structures and Algorithms	2	c0278a	15/9
c0278a	Introduction to Database Management Systems	1	c1119	5/9
c0278a	Introduction to Database Management Systems	1	c0278a	7/9
c0278a	Introduction to Database Management Systems	2	c0278a	15/9



$\sigma_{\text{Course.course_id}=\text{Tutorial.course_id}}$ (Course × Tutorial)

Course.course_id	Course.name	Tutorial.tutorial_id	Tutorial.course_id	Tutorial.date
c1119	Data Structures and Algorithms	1	c1119	5/9
c0278a	Introduction to Database Management Systems	1	c0278a	7/9
c0278a	Introduction to Database Management Systems	2	c0278a	15/9

Cartesian product

$\pi_{\text{Tutorial.date}}$

$\sigma_{\text{Course.name} = \text{"Introduction to Database Management Systems"}}$ (
 $\sigma_{\text{Course.course_id} = \text{Tutorial.course_id}}$ (**Course × Tutorial**)
)
)

Tutorial.date
7/9
15/9



$\sigma_{\text{Course.name} = \text{"Introduction to Database Management Systems"}}$ ($\sigma_{\text{Course.course_id} = \text{Tutorial.course_id}}$ (**Course × Tutorial**))

Course.course_id	Course.name	Tutorial.tutorial_id	Tutorial.course_id	Tutorial.date
c0278a	Introduction to Database Management Systems	1	c0278a	7/9
c0278a	Introduction to Database Management Systems	2	c0278a	15/9



$\sigma_{\text{Course.course_id} = \text{Tutorial.course_id}}$ (**Course × Tutorial**)

Course.course_id	Course.name	Tutorial.tutorial_id	Tutorial.course_id	Tutorial.date
c1119	Data Structures and Algorithms	1	c1119	5/9
c0278a	Introduction to Database Management Systems	1	c0278a	7/9
c0278a	Introduction to Database Management Systems	2	c0278a	15/9

Rename



Notation: $\rho_X(E)$

- Rename operator allows us to name and refer to the results of relational-algebra expressions



$\rho_X(E)$ returns the expression E under the name X

```
SELECT Tutorial.date  
FROM Course, Tutorial  
WHERE  
Coruse.name="Introduction to Database  
Management Systems" AND  
Course.course_id = Tutorial.course_id;
```

SQL

$$\pi_{\text{Tutorial.date}} \left(\sigma_{\text{Course.name}=\text{"Introduction to Database Management Systems"} \wedge \text{Course.course_id}=\text{Tutorial.course_id}} (\text{Course} \times \text{Tutorial}) \right)$$

Relational algebra (Without rename)

```
SELECT T.date  
FROM Course C, Tutorial T  
WHERE  
C.name="Introduction to Database  
Management Systems" AND  
C.course_id = T.course_id;
```

SQL

$$\pi_{T.date} \left(\sigma_{C.name=\text{"Introduction to Database Management Systems"} \wedge C.course_id=T.course_id} (\rho_C(\text{Course}) \times \rho_T(\text{Tutorial})) \right)$$

Relational algebra (With Course rename to C, Tutorial rename to T)

Section 2

Exercises

Question 1

- Given the following relational schema:
 - Student** (UID, name, age).
 - Course** (CID, title).
 - Enroll** (UID, CID) with **UID** referencing **Student** and **CID** referencing **Course**.

***UID** and **CID**, **age** are **integer**; **name** and **title** are **varchar**.

- Which of the following is (are) valid Relational Algebra expression(s)?

- $\pi_{\text{UID}}(\text{Student}) - \pi_{\text{CID}}(\text{Course})$ ✓
- $\text{Course} - \pi_{\text{UID}}(\text{Enroll})$ ✗

The left and right parts of Set difference have to be **comparable** (same number of attributes).



Question 1

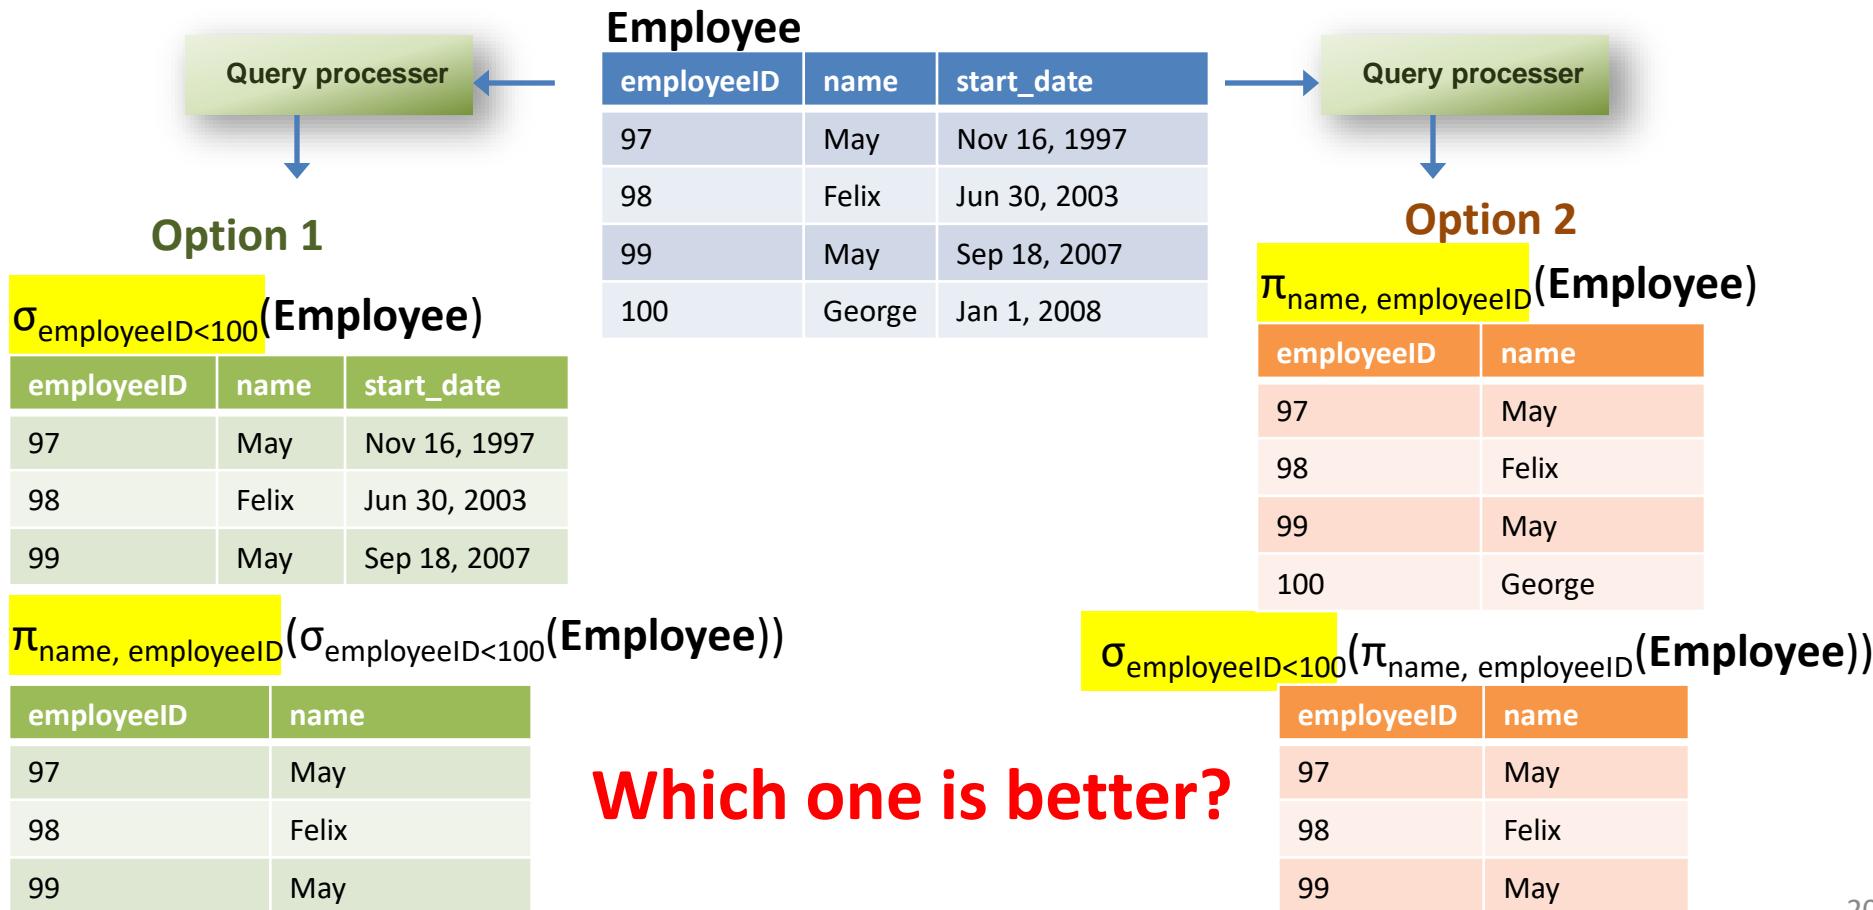
- Given the following relational schema:
 - Student (UID, name, age).
 - Course (CID, title).
 - Enroll (UID, CID) with **UID** referencing **Student** and **CID** referencing **Course**.
- ***UID** and **CID**, **age** are **integer**; **name** and **title** are **varchar**.
- Which of the following is (are) valid Relational Algebra expression(s)?
 - $\sigma_{\text{age} < 18}(\text{Student} \cup \text{Course})$  Student and Course have different attributes.
 - $\sigma_{\text{age} < 18}(\pi_{\text{UID}, \text{name}}(\text{Student}))$  No attribute “age” for selection.

Question 2

```
SELECT employee_id, name
FROM Employee
WHERE employee_id < 100;
```

SQL

- Find the employeeIDs and names of employees whose employeeID < 100.



Question 3

- Find the dept. id(s) where employees named Smith work.

Option 1

$$\pi_{W.department_id} (\sigma_{E.name='Smith'} (\sigma_{E.employee_id=W.employee_id} (\rho_E(\text{Employees}) \times \rho_W(\text{Works_in}))))$$


Now we compute the Cartesian product between **Employees** and **Works_in** first, which creates an **intermediate relation** with $10,000 * 5 = 50,000$ tuples! Can we reduce the size of intermediate relation?

Employees

employee_id	name	salary
1	Jones	26000
2	Smith	28000
...
10000

Employees (10000 tuples)

Works_in

employee_id	department_id	since
1	1	2001-1-1
2	1	2002-4-1
2	2	2005-2-2
3	3	2003-1-1
4	3	2005-1-1

Works_in (5 tuples)

```
SELECT W.department_id  
FROM Employees E, Works_in W  
WHERE E.name = 'Smith' AND  
E.employee_id = W.employee_id;
```

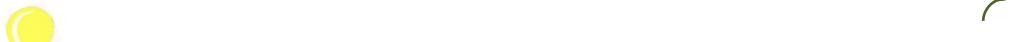
Question 3

- Find the dept. id(s) where employees named Smith work.

Option 1

$$\pi_{W.department_id} (\sigma_{E.name='Smith'} (\sigma_{E.employee_id=W.employee_id} (\rho_E(Employees) \times \rho_W(Works_in))))$$

Option 2

$$\pi_{W.department_id} (\sigma_{E.employee_id=W.employee_id} (\sigma_{E.name='Smith'} (\rho_E(Employees)) \times \rho_W(Works_in)))$$


We apply selection on Employees before the Cartesian product. If there is only two employee named Smith, the **intermediate relation** can be reduced to $2 \times 5 = 10$ tuples!

Employees

employee_id	name	salary
1	Jones	26000
2	Smith	28000
...
10000

Employees (10000 tuples)

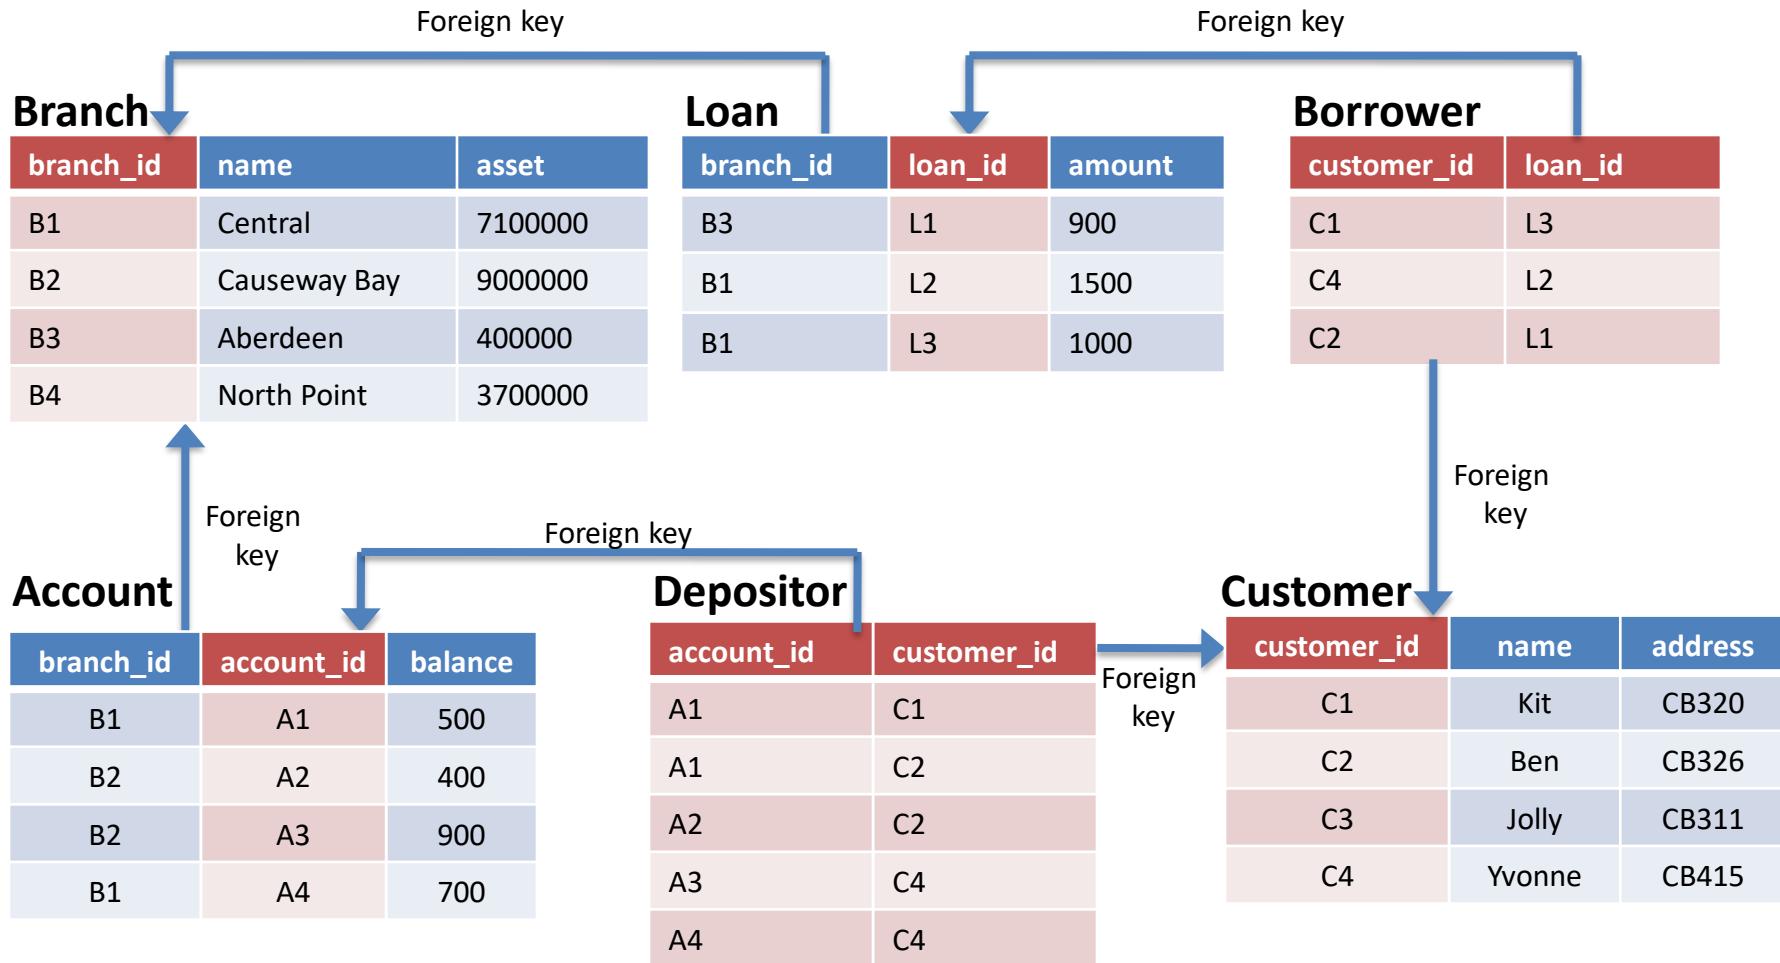
Works_in

employee_id	department_id	since
1	1	2001-1-1
2	1	2002-4-1
2	2	2005-2-2
3	3	2003-1-1
4	3	2005-1-1

Works_in (5 tuples)

```
SELECT W.department_id  
FROM Employees E, Works_in W  
WHERE E.name = 'Smith' AND  
      E.employee_id = W.employee_id;
```

Banking example



Question 4

- Find the names of all customers who have a loan, an account, or both in a bank.

```
SELECT customer_id  
FROM Borrower  
UNION  
SELECT customer_id  
FROM Depositor
```

$$\pi_{\text{customer_id}}(\text{Borrower}) \cup \pi_{\text{customer_id}}(\text{Depositor})$$

- Find the names of all customers who have both a loan and an account in a bank.

```
SELECT customer_id  
FROM Borrower  
INTERSECT  
SELECT customer_id  
FROM Depositor
```

$$\pi_{\text{customer_id}}(\text{Borrower}) \cap \pi_{\text{customer_id}}(\text{Depositor})$$

Wait! Do we have set intersection in relational algebra ☺?

$$\pi_{\text{customer_id}}(\text{Borrower}) - (\pi_{\text{customer_id}}(\text{Borrower}) - \pi_{\text{customer_id}}(\text{Depositor}))$$


Additional operators

- These operations do not add any power to relational algebra, but simplify common queries.
- Set intersection (\cap)
- Natural join (\bowtie)
- Assignment (\leftarrow)
- Left outer join (\bowtie_l), Right outer join (\bowtie_r)
- Division (\div)

Section 3

Additional operators

Motivation

- The fundamental operators of the relational algebra introduced are sufficient to express any relational-algebra query.
- However, if we restrict ourselves to just the fundamental operators, certain common queries are tedious to express.
- Therefore, we define additional operators that do not add any power to the algebra, but simplify common queries

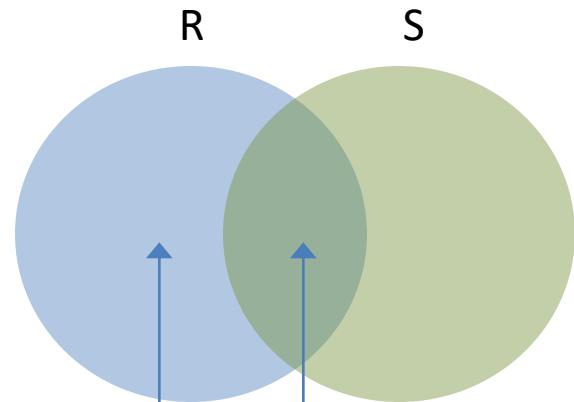
That is to say, for each additional operator, we can give an equivalent expression that use only the fundamental operators.



Additional operators

- Set intersection (\cap)
- Natural join (\bowtie)
- Assignment (\leftarrow)
- Left outer join (\bowtie_l), Right outer join (\bowtie_r)
- Division (\div)

Set intersection



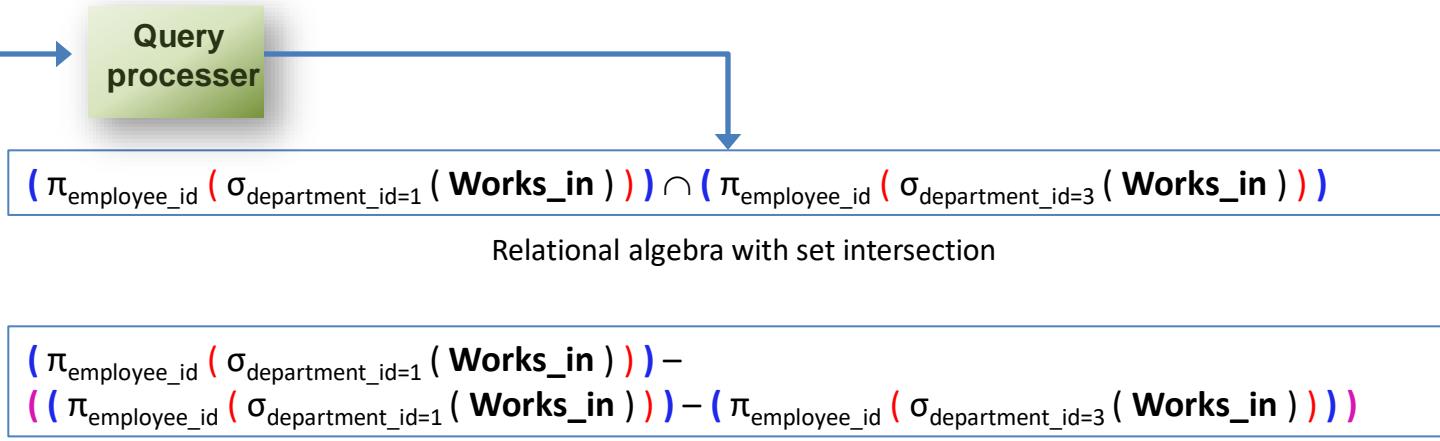
- $R \cap S = R - (R - S)$

- R and S must have the **same number of attributes** and **attribute data types are compatible**.

Example

- **Query:** Find the employee_id of employees who work in department 1 **and** department 3.

```
SELECT employee_id  
FROM Works_in  
WHERE department_id=1  
  
INTERSECT  
  
SELECT employee_id  
FROM Works_in  
WHERE department_id=3
```



Set intersection

For your reference,
there is another way to
answer the same query
by joining the
Works_in table with
itself.



```
SELECT DISTINCT employee_id  
FROM Works_in W1, Works_in W2  
WHERE  
W1.employee_id = W2.employee_id AND  
W1.department_id=1 AND  
W2. department_id = 3
```

SQL

Query
processor

Works_in

employee_id	department_id
1	1
2	1
1	3

$$\pi_{\text{employee_id}} (\sigma_{W1.\text{department_id}=1 \wedge W2.\text{department_id}=3} (\sigma_{W1.\text{employee_id}=W2.\text{employee_id}} (\rho_{W1}(\text{Works_in}) \times \rho_{W2}(\text{Works_in}))))$$

Set intersection

For your reference,
there is another way to
answer the same query
by joining the
Works_in table with
itself.



SQL

```
SELECT DISTINCT employee_id
FROM Works_in W1, Works_in W2
WHERE
W1.employee_id = W2.employee_id AND
W1.department_id=1 AND
W2. department_id = 3
```

$\rho_{W1}(\text{Works_in}) \times \rho_{W2}(\text{Works_in})$

W1.employee_id	W1.department_id	W2.employee_id	W2.department_id
1	1	1	1
1	1	2	1
1	1	1	3
2	1	1	1
2	1	2	1
2	1	1	3
1	3	1	1
1	3	2	1
1	3	1	3

Query
processor

↑

Works_in	
employee_id	department_id
1	1
2	1
1	3

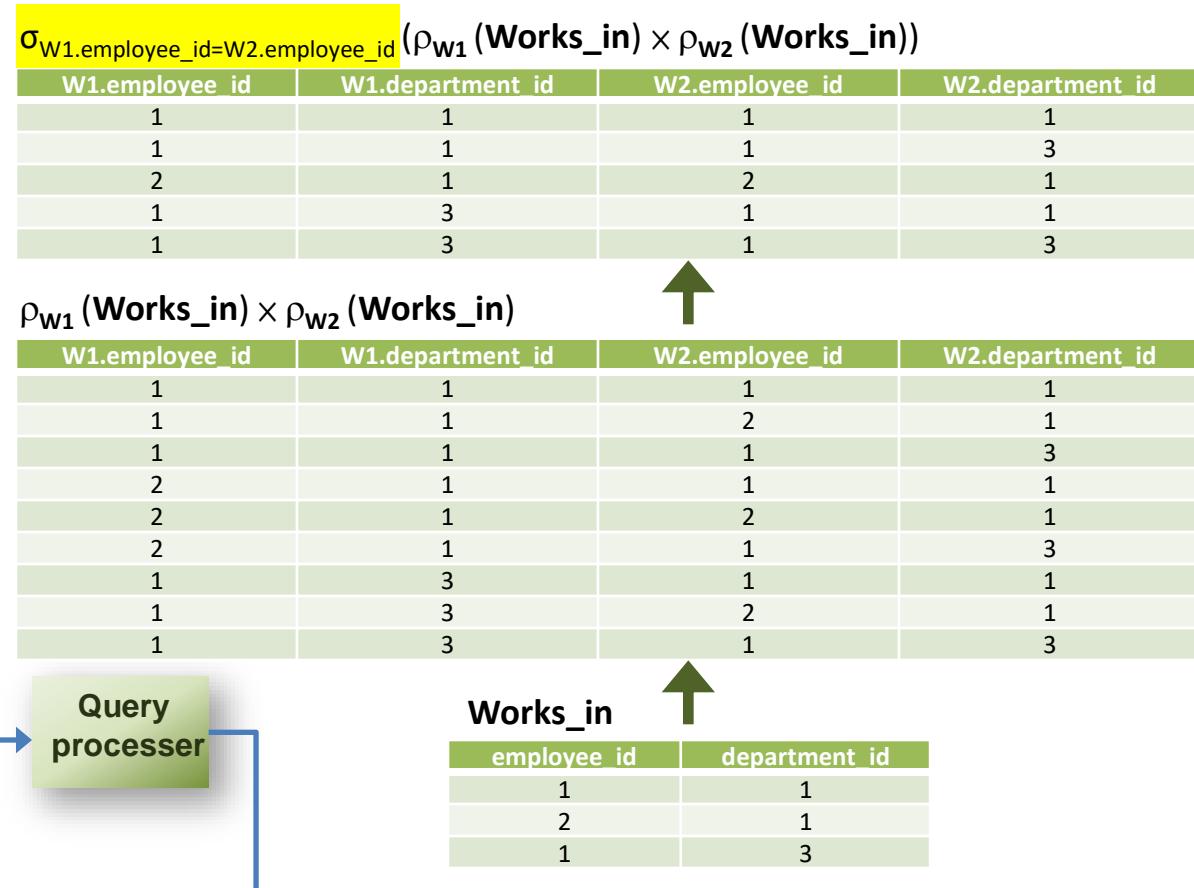
$\pi_{\text{employee_id}} (\sigma_{W1.\text{department_id}=1 \wedge W2.\text{department_id}=3} (\sigma_{W1.\text{employee_id}=W2.\text{employee_id}} (\rho_{W1}(\text{Works_in}) \times \rho_{W2}(\text{Works_in}))))$

Set intersection

For your reference,
there is another way to
answer the same query
by joining the
Works_in table with
itself.



```
SELECT DISTINCT employee_id
FROM Works_in W1, Works_in W2
WHERE
W1.employee_id = W2.employee_id AND
W1.department_id=1 AND
W2. department_id = 3
SQL
```



$\pi_{employee_id} (\sigma_{W1.department_id=1 \wedge W2.department_id=3} (\sigma_{W1.employee_id=W2.employee_id} (\rho_{W1}(\text{Works_in}) \times \rho_{W2}(\text{Works_in}))))$

Set intersection

$$\sigma_{W1.department_id=1 \wedge W2.department_id=3} (\sigma_{W1.employee_id=W2.employee_id} (\rho_{W1}(\text{Works_in}) \times \rho_{W2}(\text{Works_in})))$$

W1.employee_id	W1.department_id	W2.employee_id	W2.department_id
1	1	1	3

For your reference,
there is another way to
answer the same query
by joining the
Works_in table with
itself.



```
SELECT DISTINCT employee_id
FROM Works_in W1, Works_in W2
WHERE
W1.employee_id = W2.employee_id AND
W1.department_id=1 AND
W2.department_id = 3
```

SQL

$$\sigma_{W1.employee_id=W2.employee_id} (\rho_{W1}(\text{Works_in}) \times \rho_{W2}(\text{Works_in}))$$

W1.employee_id	W1.department_id	W2.employee_id	W2.department_id
1	1	1	1
1	1	1	3
2	1	2	1
1	3	1	1
1	3	1	3

$$\rho_{W1}(\text{Works_in}) \times \rho_{W2}(\text{Works_in})$$

W1.employee_id	W1.department_id	W2.employee_id	W2.department_id
1	1	1	1
1	1	2	1
1	1	1	3
2	1	1	1
2	1	2	1
2	1	1	3
1	3	1	1
1	3	2	1
1	3	1	3

Query
processor

Works_in

employee_id	department_id
1	1
2	1
1	3

$$\pi_{\text{employee_id}} (\sigma_{W1.department_id=1 \wedge W2.department_id=3} (\sigma_{W1.employee_id=W2.employee_id} (\rho_{W1}(\text{Works_in}) \times \rho_{W2}(\text{Works_in}))))$$

Natural join

- Usually, a query that involves a Cartesian product includes a selection operation on the result of the Cartesian product.
 - The selection operation most often requires that all **attributes that are common** to the relations that are involved in the Cartesian product **be equated**.

$$r \bowtie s = \pi_{R \cup S} (\sigma_{\underbrace{r.A_1 = s.A_1 \wedge r.A_2 = s.A_2 \wedge \dots \wedge r.A_n = s.A_n}_{\text{Requires the common attributes to be equated}}} (r \times s))$$

Requires the common attributes to be equated

- where $R \cap S = \{A_1, A_2, \dots, A_n\}$

Commutative: $r \bowtie s = s \bowtie r$

Natural join

- The schema of $R \bowtie S$ is R-schema \cup S-schema
(repeated attributes are removed)
 - For each pair of tuples t_r from R and t_s from S,
 - If t_r and t_s share the same value over each of the common attributes in R and S,
 - Tuple t will be added to the result of $R \bowtie S$.

Natural join

R S

A	B	A	C
1	1	1	2
1	2	2	1
2	3		

- Common attributes: $R \cap S = \{A\}$

- Attributes of the resulting relation: $R \cup S = \{A, B, C\}$

R \times S

R.A	R.B	S.A	S.C
1	1	1	2
1	1	2	1
1	2	1	2
1	2	2	1
2	3	1	2
2	3	2	1



$\sigma_{R.A=S.A}(R \times S)$

R.A	R.B	S.A	S.C
1	1	1	2
1	2	1	2
2	3	2	1



$\pi_{A,B,C}(\sigma_{R.A=S.A}(R \times S))$

A	B	C
1	1	2
1	2	2
2	3	1

equivalent to:

R \bowtie S

A	B	C
1	1	2
1	2	2
2	3	1

Assignment

- It is convenient to write a relational-algebra expression by assigning parts of it to **temporary relation variables**.
- The assignment operator, denoted by “ \leftarrow ”, works like assignment operator “=” in programming language.
 - temp1 $\leftarrow \pi_a(R)$
 - temp2 $\leftarrow \pi_a(S)$
 - result $\leftarrow \text{temp1} - \text{temp2}$

With the assignment operator, a query can be written as a sequential program.

temp1, temp2, result are called “**relation variable**”.



Outer join

- The outer-join operator is an extension of the join operation to deal with missing information.

Customer

customer_id	name	address
C1	Kit	CB320
C2	Ben	CB326
C3	Jolly	CB311
C4	Yvonne	CB415

Depositor

account_id	customer_id
A1	C1
A1	C2
A2	C2
A3	C4
A4	C4

Customer \bowtie Depositor

customer_id	name	address	account_id
C1	Kit	CB320	A1
C2	Ben	CB326	A1
C2	Ben	CB326	A2
C4	Yvonne	CB415	A3
C4	Yvonne	CB415	A4

Natural join (e.g., joining Customer and Depositor) result in a table **without the information of customers (e.g., C3 in this case) who has no account**.

Outer join

- Natural join result, plus
 - The **tuples that do not match any tuples from the other side.**



Left outer join

$$R \Join S = (R \bowtie S) \cup (R - \pi_R(R \bowtie S)) \times \{ (null, \dots, null) \}$$

Customer

customer_id	name	address
C1	Kit	CB320
C2	Ben	CB326
C3	Jolly	CB311
C4	Yvonne	CB415

Depositor

account_id	customer_id
A1	C1
A1	C2
A2	C2
A3	C4
A4	C4

Customer \bowtie Depositor

customer_id	name	address	account_id
C1	Kit	CB320	A1
C2	Ben	CB326	A1
C2	Ben	CB326	A2
C4	Yvonne	CB415	A3
C4	Yvonne	CB415	A4

Let's illustrate why the left outer join is defined as
 $(R \bowtie S) \cup (R - \pi_R(R \bowtie S)) \times \{ (null, \dots, null) \}$
through a step-by-step illustration.



Left outer join

$$R \text{ } \bowtie \text{ } S = (R \bowtie S) \cup (R - \pi_R(R \bowtie S)) \times \{ (null, \dots, null) \}$$

Customer

customer_id	name	address
C1	Kit	CB320
C2	Ben	CB326
C3	Jolly	CB311
C4	Yvonne	CB415

Depositor

account_id	customer_id
A1	C1
A1	C2
A2	C2
A3	C4
A4	C4

Customer \bowtie Depositor

customer_id	name	address	account_id
C1	Kit	CB320	A1
C2	Ben	CB326	A1
C2	Ben	CB326	A2
C4	Yvonne	CB415	A3
C4	Yvonne	CB415	A4

Customer - $\pi_{\text{Customer's attributes}}$ (Customer \bowtie Depositor)

customer_id	name	address
C3	Jolly	CB311

Finding missed tuples in the natural join

$R - \pi_R(R \bowtie S)$ is to generate the tuples in R (i.e., Customer) that are missed in the natural join.

i.e., C3 Jolly in Customer doesn't have any matched records in Depositor, we use this part to recover Jolly's record.



Left outer join

$$R \text{ } \bowtie \text{ } S = (R \bowtie S) \cup (R - \pi_R(R \bowtie S)) \times \{ (null, \dots, null) \}$$

Customer

customer_id	name	address
C1	Kit	CB320
C2	Ben	CB326
C3	Jolly	CB311
C4	Yvonne	CB415

Depositor

account_id	customer_id
A1	C1
A1	C2
A2	C2
A3	C4
A4	C4

Customer \bowtie Depositor

customer_id	name	address	account_id
C1	Kit	CB320	A1
C2	Ben	CB326	A1
C2	Ben	CB326	A2
C4	Yvonne	CB415	A3
C4	Yvonne	CB415	A4

Customer - $\pi_{\text{Customer's attributes}}$ (Customer \bowtie Depositor)

customer_id	name	address
C3	Jolly	CB311

Customer - $\pi_{\text{Customer's attributes}}$ (Customer \bowtie Depositor) $\times \{ (null, \dots, null) \}$

customer_id	name	address	account_id
C3	Jolly	CB311	null

Constructing the missed tuple by adding null value to extra attributes

The “ $\times \{ (null, \dots, null) \}$ ” part simply add back the remaining column values as **null** because there is no matched records in S (i.e., Depositor).



Left outer join

$$R \bowtie S = (R \bowtie S) \cup (R - \pi_R(R \bowtie S)) \times \{ (null, \dots, null) \}$$

Customer

customer_id	name	address
C1	Kit	CB320
C2	Ben	CB326
C3	Jolly	CB311
C4	Yvonne	CB415

Depositor

account_id	customer_id
A1	C1
A1	C2
A2	C2
A3	C4
A4	C4

Customer \bowtie Depositor

customer_id	name	address	account_id
C1	Kit	CB320	A1
C2	Ben	CB326	A1
C2	Ben	CB326	A2
C4	Yvonne	CB415	A3
C4	Yvonne	CB415	A4

Customer - $\pi_{\text{Customer's attributes}}$ (Customer \bowtie Depositor)

customer_id	name	address
C3	Jolly	CB311

Customer - $\pi_{\text{Customer's attributes}}$ (Customer \bowtie Depositor) $\times \{ (null, \dots, null) \}$

customer_id	name	address	account_id
C3	Jolly	CB311	null

(Customer \bowtie Depositor) \cup (Customer - $\pi_{\text{Customer's attributes}}$ (Customer \bowtie Depositor)) $\times \{ (null, \dots, null) \}$

customer_id	name	address	account_id
C1	Kit	CB320	A1
C2	Ben	CB326	A1
C2	Ben	CB326	A2
C3	Jolly	CB311	null
C4	Yvonne	CB415	A3
C4	Yvonne	CB415	A4

equivalent to: Customer \bowtie Depositor

Right outer join

$$R \bowtie S = (R \bowtie S) \cup (S - \pi_S(R \bowtie S)) \times \{ (null, \dots, null) \}$$

Customer

customer_id	name	address
C1	Kit	CB320
C2	Ben	CB326
C3	Jolly	CB311
C4	Yvonne	CB415

Depositor

account_id	customer_id
A1	C1
A1	C2
A2	C2
A3	C4
A4	C5

Customer \bowtie Depositor

customer_id	name	address	account_id
C1	Kit	CB320	A1
C2	Ben	CB326	A1
C2	Ben	CB326	A2
C4	Yvonne	CB415	A3

Let's illustrate why the right outer join is defined as

$(R \bowtie S) \cup (S - \pi_S(R \bowtie S)) \times \{ (null, \dots, null) \}$

through a step-by-step illustration.



Right outer join

$$R \bowtie S = (R \bowtie S) \cup (S - \pi_S(R \bowtie S)) \times \{ (null, \dots, null) \}$$

Customer

customer_id	name	address
C1	Kit	CB320
C2	Ben	CB326
C3	Jolly	CB311
C4	Yvonne	CB415

Depositor

account_id	customer_id
A1	C1
A1	C2
A2	C2
A3	C4
A4	C5

Customer \bowtie Depositor

customer_id	name	address	account_id
C1	Kit	CB320	A1
C2	Ben	CB326	A1
C2	Ben	CB326	A2
C4	Yvonne	CB415	A3

Depositor - $\pi_{\text{Depositor's attributes}}$ (Customer \bowtie Depositor)

customer_id	account_id
C5	A4

Finding missed tuples in the natural join

$S - \pi_S(R \bowtie S)$ is to generate the tuples in S (i.e., Depositor) that are missed in the natural join.
i.e., C5 in Depositor doesn't have any matched records in Customer, we use this part to recover C5's record.



Right outer join

$$R \bowtie S = (R \bowtie S) \cup (S - \pi_S(R \bowtie S)) \times \{(\text{null}, \dots, \text{null})\}$$

Customer

customer_id	name	address
C1	Kit	CB320
C2	Ben	CB326
C3	Jolly	CB311
C4	Yvonne	CB415

Depositor

account_id	customer_id
A1	C1
A1	C2
A2	C2
A3	C4
A4	C5

Customer \bowtie Depositor

customer_id	name	address	account_id
C1	Kit	CB320	A1
C2	Ben	CB326	A1
C2	Ben	CB326	A2
C4	Yvonne	CB415	A3

Depositor - $\pi_{\text{Depositor's attributes}}$ (Customer \bowtie Depositor)

customer_id	account_id
C5	A4

Depositor - $\pi_{\text{Depositor's attributes}}$ (Customer \bowtie Depositor) $\times \{(\text{null}, \dots, \text{null})\}$

customer_id	name	address	account_id
C5	null	null	A4

Constructing the missed tuple by adding *null* value to extra attributes

The “ $\times \{(\text{null}, \dots, \text{null})\}$ ” part simply add back the remaining column values as *null* because there is no matched records in R (i.e., Customer).



Right outer join

$$R \bowtie S = (R \bowtie S) \cup (S - \pi_S(R \bowtie S)) \times \{ (null, \dots, null) \}$$

Customer

customer_id	name	address
C1	Kit	CB320
C2	Ben	CB326
C3	Jolly	CB311
C4	Yvonne	CB415

Depositor

account_id	customer_id
A1	C1
A1	C2
A2	C2
A3	C4
A4	C5

Customer \bowtie Depositor

customer_id	name	address	account_id
C1	Kit	CB320	A1
C2	Ben	CB326	A1
C2	Ben	CB326	A2
C4	Yvonne	CB415	A3

Depositor - $\pi_{\text{Depositor's attributes}}$ (Customer \bowtie Depositor)

customer_id	account_id
C5	A4

Depositor - $\pi_{\text{Depositor's attributes}}$ (Customer \bowtie Depositor) $\times \{ (null, \dots, null) \}$

customer_id	name	address	account_id
C5	null	null	A4

(Customer \bowtie Depositor) \cup (Depositor - $\pi_{\text{Depositor's attributes}}$ (Customer \bowtie Depositor)) $\times \{ (null, \dots, null) \}$

customer_id	name	address	account_id
C1	Kit	CB320	A1
C2	Ben	CB326	A1
C2	Ben	CB326	A2
C4	Yvonne	CB415	A3
C5	null	null	A4

equivalent to: Customer \bowtie Depositor

Division

Notation: $R \div S$

Definition

Let $S \subseteq R$

$R \div S = \{ t \mid t \in \pi_{R-S}(R) \wedge (\forall s \in S, (t \cup s) \in R) \}$

Condition 1. A resulting tuple t has to be in the relation $\pi_{R-S}(R)$

A
1
2
3
4

The attributes in relation S is a subset of the attributes in relation R .

Condition 2. And if we combine t with each tuple $s \in S$, all the combined tuples have to be included in R .

1 For each tuple $t \in \pi_{R-S}(R)$

2 “ $\forall s \in S, t \cup s$ ” part:
Generate tuples by union t with all tuples $s \in S$

3 “ $\forall s \in S, (t \cup s) \in R$ ” part:
Then we check if both tuples generated are in R .

$t_1 \in \pi_{R-S}(R)$

A
1

$\forall s \in S, (t_1 \cup s)$

A	B
1	1
1	2

In this case, not both tuples are in R , so t_1 is **NOT** in result of $R \div S$

$t_2 \in \pi_{R-S}(R)$

A
2

$\forall s \in S, (t_2 \cup s)$

A	B
2	1
2	2

In this case, both tuples are in R , so t_2 is in result of $R \div S$

R

A	B
1	1
2	1
2	2
3	3
4	1
4	2
4	3

S

B
1
2

$R \div S$

A
2
4

Division

Observation

Let's focus on the result of $R \div S$ (say, the tuple A=2), it means that

- For the tuples with values in attribute A equals to 2 in relation R,
- Those tuples's values in attribute B covers **ALL** values in attribute B of S.

- Division is used to express queries with “all”

- Find the IDs of all students who have taken **all** CS courses (dpt_id = 1).



Student

student_id	name	dpt_id
1	Peter	1
2	Sharon	1
3	David	2
4	Joe	3

Takes

student_id	course_id	Grade
1	1	A
1	2	B
1	3	A+
2	3	B-
3	3	B
4	1	C
4	2	A-

Course

course_id	title	dpt_id	credit
1	Intro to DB	1	6
2	Programming I	1	6
3	Accounting	2	6

Division



Step 1. All part (the relation S): What is the course ID of all CS courses ($dpt_id = 1$)?

course_id
1
2

$\pi_{course_id}(\sigma_{dpt_id=1}(Course))$



$\sigma_{dpt_id=1}(Course)$

course_id	title	dpt_id	credit
1	Intro to DB	1	6
2	Programming I	1	6



Student

student_id	name	dpt_id
1	Peter	1
2	Sharon	1
3	David	2
4	Joe	3

Takes

student_id	course_id	Grade
1	1	A
1	2	B
1	3	A+
2	3	B-
3	3	B
4	1	C
4	2	A-

Course

course_id	title	dpt_id	credit
1	Intro to DB	1	6
2	Programming I	1	6
3	Accounting	2	6

Division



Step 2. Dividend part (the relation R): What is the list of Takes tuples (i.e., all $\langle \text{student_id}, \text{course_id} \rangle$ pairs)?

$\pi_{\text{student_id}, \text{course_id}}(\text{Takes})$

student_id	course_id
1	1
1	2
1	3
2	3
3	3
4	1
4	2

$\pi_{\text{course_id}}(\sigma_{\text{dpt_id}=1}(\text{Course}))$

course_id
1
2



Student

student_id	name	dpt_id
1	Peter	1
2	Sharon	1
3	David	2
4	Joe	3

Takes

student_id	course_id	Grade
1	1	A
1	2	B
1	3	A+
2	3	B-
3	3	B
4	1	C
4	2	A-

Course

course_id	title	dpt_id	credit
1	Intro to DB	1	6
2	Programming I	1	6
3	Accounting	2	6

Division



Step 3. Division: Which student in $\pi_{\text{student_id}, \text{course_id}}$ (Takes) takes **all** CS courses $\pi_{\text{course_id}} (\sigma_{\text{dpt_id}=1}(\text{course}))$?.

$\pi_{\text{student_id}, \text{course_id}}$ (Takes)

$\pi_{\text{course_id}} (\sigma_{\text{dpt_id}=1}(\text{Course}))$

$\pi_{\text{student_id}, \text{course_id}} (\text{Takes}) \div \pi_{\text{course_id}} (\sigma_{\text{dpt_id}=1}(\text{Course}))$

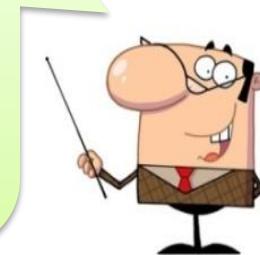
student_id	course_id
1	1
1	2
1	3
2	3
3	3
4	1
4	2

course_id
1
2

student_id
1
4

Explanation: Let's focus on $\text{student_id} = 1$ in the result

- It means that, for the tuples in **Takes** with $\text{student_id}=1$,
- Those tuple's value in course_id attribute covers **all** 1 and 2 (i.e., the CS courses).
- That is to say, student with student_id 1 takes **ALL** CS courses.
- The same argument applies to student_id 4.



Division



Division has a property: $R \times S \div S = R$

R	
A	B
1	1
1	2
2	3

S	
C	D
1	2
2	2

$R \times S$

A	B	C	D
1	1	1	2
1	1	2	2
1	2	1	2
1	2	2	2
2	3	1	2
2	3	2	2

$(R \times S) \div S$

A	B
1	1
1	2
2	3

= R

An intuitive property of the division operator of the relational algebra is simply that it is the inverse of the cartesian product. For example, if you have two relations R and S, then, if U is a relation defined as the cartesian product of them: $U = R \times S$, the division is the operator such that $U \div R = S$ and $U \div S = R$

Section 4

Extended operators

Aggregation

- Aggregation function takes a collection of values and returns a single-valued result.
 - e.g., avg, min, max, sum, count, count-distinct
- Aggregate operation in relational algebra:
 - **Grouping** - Divides the tuples into groups.
 - **Aggregation** - Computes an aggregation function in each group to create a result tuple.

Aggregation

$\text{g}_{G_1, G_2, \dots, G_n} F_1(A_1), F_2(A_2), \dots, F_n(A_n) (E)$

- **E** – a relation (can be a result of relational algebra expression).
- **G₁, ..., G_n** – attributes used to form groups.
 - Tuples with the same values in **G₁** to **G_n** are put into the same group.
 - **G** can be empty, which means that the whole relation is one group.
- **F_i(A_i)** – an aggregate function applied on an attribute.

Aggregation

Account

branch_id	account_id	balance
B1	A1	500
B2	A2	400
B2	A3	900
B1	A4	700

```
 $\rho_{\text{Result}(\text{branch\_id}, \text{sum\_of\_balance})} ($   
     $\text{branch\_id } g \text{ sum(balance)} (\text{Account})$   
 $)$ 
```

Result



branch_id	sum_of_balance
B1	1200
B2	1300

- Step 1. Let's group the tuples in Account according to their branch_id.
- Step 2. Then aggregate the tuples in each group by summing their values in the balance attribute.
- Step 3. Since the resulting relation has no name after aggregation, we use renaming operator to give name to the relation and attributes.



Aggregation

- Note that grouping **can be done on multiple attributes.**
- E.g., in this case, we group tuples in Student with the same values in both dpt_id and gender attributes.
- i.e., We are finding the number of male / female students in each department.



Student

student_id	name	dpt_id	gender
1	Peter	1	M
2	Sharon	1	F
3	David	2	M
4	Joe	2	M
5	Betty	1	F

$$\rho_{\text{Result}(\text{dpt_id}, \text{gender}, \text{count})} ($$

$$g_{\text{count}(\text{Student})} (\text{dpt_id}, \text{gender}))$$


Result

dpt_id	gender	count
1	M	1
1	F	2
2	M	2

Section 5

Algebraic properties

Transformation of expression

- A query can be expressed in several different ways, with different cost of evaluation.
- Two relational-algebra expressions are said to be equivalent if, on every legal database instance, the two expressions generate the same set of tuples.

In the following discussions, we treat a relation as a set of tuples, the order of the tuples is irrelevant.



Equivalence rules

- Rule 1. Only the final operations in a sequence of projection operations are needed.

$$\pi_{L1} (\pi_{L2} (\dots (\pi_{Ln} (E)) \dots)) = \pi_{L1} (E)$$

$\pi_{\text{employee_id}} (\pi_{\text{employee_id}, \text{salary}} (\text{Employee}))$

employee_id
1
2
3

$\pi_{\text{employee_id}, \text{salary}} (\text{Employee})$

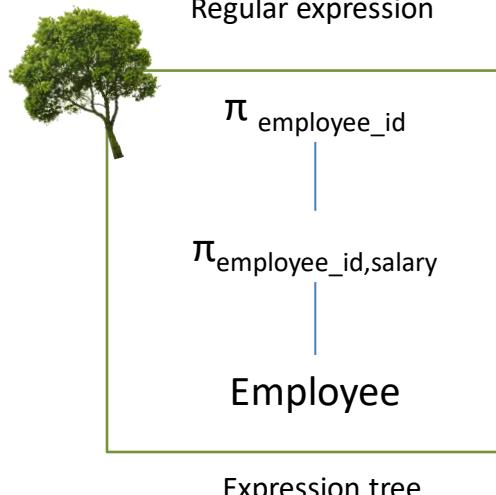
employee_id	name
1	Jones
2	Smith
3	Smith

Employee

employee_id	name	salary
1	Jones	26000
2	Smith	28000
3	Smith	24000

$\pi_{\text{employee_id}} ($
 $\pi_{\text{employee_id}, \text{salary}} (\text{Employee})$
 $)$

Regular expression



Expression tree:

- Tells which operator is executed ahead of another.
- It allows transformation of the execution order by applying the equivalence rules.
(Alter the tree)

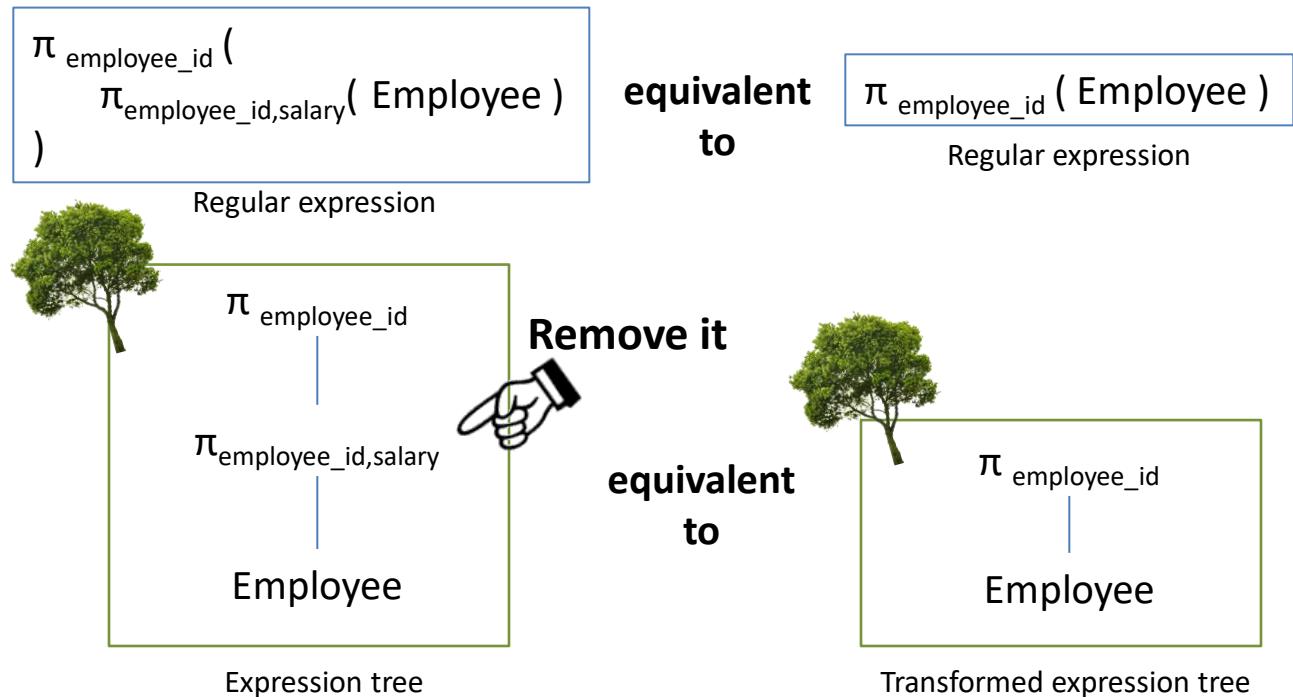


Equivalence rules

- Rule 1. Only the final operations in a sequence of projection operations are needed.

$$\pi_{L1} (\pi_{L2} (\dots (\pi_{Ln} (E)) \dots)) = \pi_{L1} (E)$$

$\pi_{\text{employee_id}}$ (Employee)		
employee_id	name	salary
1	Jones	26000
2	Smith	28000
3	Smith	24000



Equivalence rules

- Rule 2. Conjunctive selection operations can be deconstructed into a sequence of individual selections.

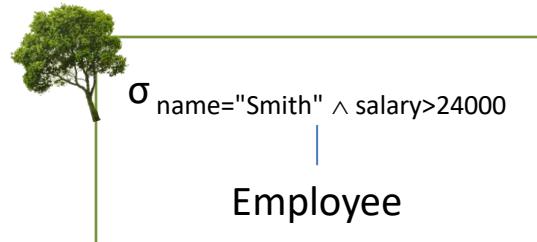
$$\sigma_{p_1 \wedge p_2}(E) = \sigma_{p_1}(\sigma_{p_2}(E))$$

$\sigma_{\text{name}=\text{"Smith"} \wedge \text{salary}>24000}(\text{Employee})$

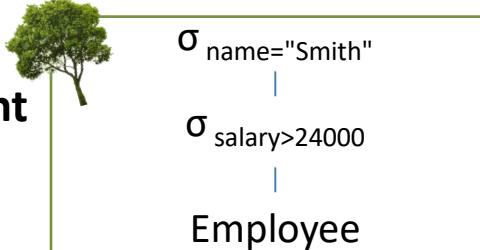
$\sigma_{\text{name}=\text{"Smith"} \wedge \text{salary}>24000}(\text{Employee})$

$\sigma_{\text{name}=\text{"Smith"} \wedge \text{salary}>24000}(\text{Employee})$

employee_id	name	salary
2	Smith	28000



equivalent
to



employee_id	name	salary
1	Jones	26000
2	Smith	28000

employee_id	name	salary
1	Jones	26000
2	Smith	28000
3	Smith	24000

You may wonder why breaking down the conjunctive selections is useful.

We will show that it is useful **to reduce temporary result**.

We can try to push each selection predicates down the tree (to perform selection as early as possible).



Equivalence rules

Rule 3. Selection operations are commutative.

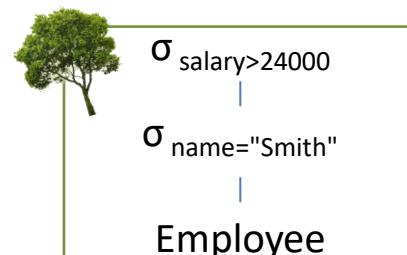
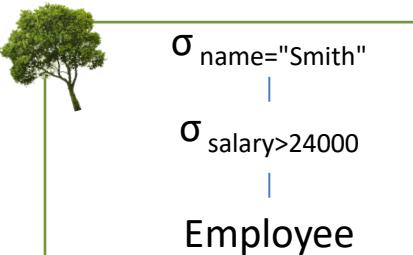
$$\sigma_{p1}(\sigma_{p2}(E)) = \sigma_{p2}(\sigma_{p1}(E))$$

$\sigma_{name='Smith'}(\sigma_{salary>24000}(Employee))$		
employee_id	name	salary
2	Smith	28000

$\sigma_{name='Smith'}($
 $\sigma_{salary>24000}(Employee)$
 $)$

$\sigma_{salary>24000}($
 $\sigma_{name='Smith'}(Employee)$
 $)$

$\sigma_{salary>24000}(Employee)$		
employee_id	name	salary
1	Jones	26000
2	Smith	28000
4	David	25000



$\sigma_{salary>24000}(\sigma_{name='Smith'}(Employee))$		
employee_id	name	salary
2	Smith	28000

$\sigma_{name='Smith'}(Employee)$		
employee_id	name	salary
2	Smith	28000
3	Smith	24000

Employee		
employee_id	name	salary
1	Jones	26000
2	Smith	28000
3	Smith	24000
4	David	25000



Note that the two executions have different costs. In particular, the size of their temporary relations are different.

employee_id	name	salary
1	Jones	26000
2	Smith	28000
3	Smith	24000
4	David	25000

Equivalence rules

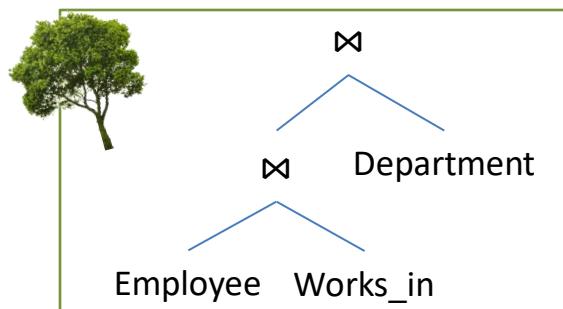
$$r \bowtie s = \pi_{R \cup S}(\sigma_{r.A_1 = s.A_1 \wedge r.A_2 = s.A_2 \wedge \dots \wedge r.A_n = s.A_n}(r \times s))$$



Rule 4. Natural join operations are associative.

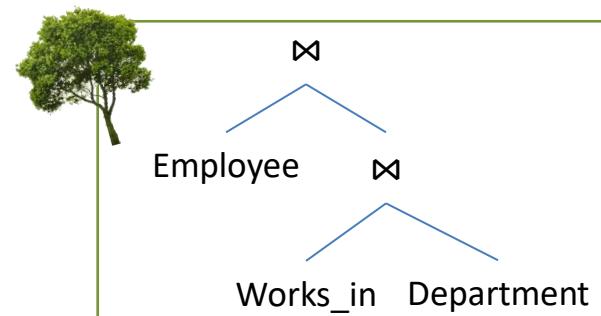
$$(E_1 \bowtie E_2) \bowtie E_3 = E_1 \bowtie (E_2 \bowtie E_3)$$

(Employee \bowtie Works_in) \bowtie Department



Employee \bowtie (Works_in \bowtie Department)

equivalent
to



Although both expression trees return the **same resulting relation**, these two expression trees have **different costs** because the size of their temporary relations are different

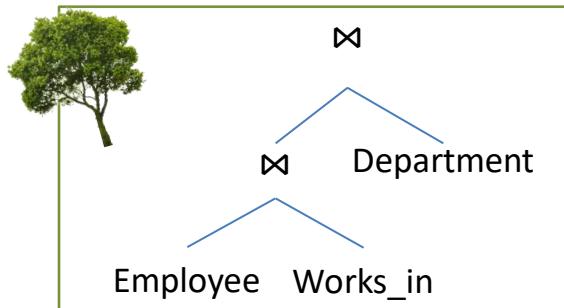


Equivalence rules



Expression tree A.

(Employee \bowtie Works_in) \bowtie Department



(Employee \bowtie Works_in) \bowtie Department

employee_id	Employee.name	salary	department_id	Department.name
1	Jones	26000	1	Toys
2	Smith	28000	1	Toys
2	Smith	28000	2	Tools
3	Parker	35000	3	Food
4	Smith	24000	3	Food

Employee \bowtie Works_in

employee_id	name	salary	department_id
1	Jones	26000	1
2	Smith	28000	1
2	Smith	28000	2
3	Parker	35000	3
4	Smith	24000	3

Natural join evaluates $4 * 5 = 20$ combinations
result temporary relation consists of 5 tuples and 4 columns.

Employee

employee_id	name	salary
1	Jones	26000
2	Smith	28000
3	Parker	35000
4	Smith	24000

Works_in

employee_id	department_id
1	1
2	1
2	2
3	3
4	3

Department

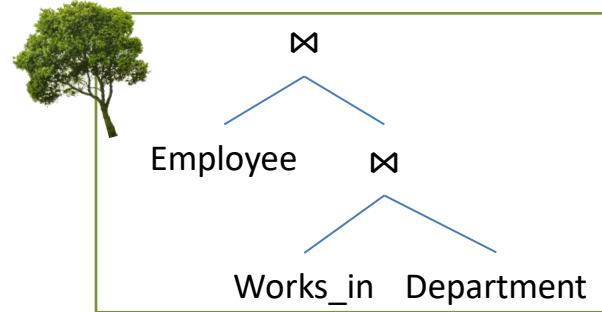
department_id	name
1	Toys
2	Tools
3	Food

Equivalence rules



Expression tree B.

Employee \bowtie (Works_in \bowtie Department)



Employee \bowtie (Works_in \bowtie Department)

employee_id	Employee.name	salary	department_id	Department.name
1	Jones	26000	1	Toys
2	Smith	28000	1	Toys
2	Smith	28000	2	Tools
3	Parker	35000	3	Food
4	Smith	24000	3	Food

Works_in \bowtie Department

employee_id	department_id	name
1	1	Toys
2	1	Toys
2	2	Tools
3	3	Food
4	3	Food

Natural join evaluates $5 * 3 = 15$ combinations
result temporary relation consists of 5 tuples and 3 columns.

Department

department_id	name
1	Toys
2	Tools
3	Food

Employee		
employee_id	name	salary
1	Jones	26000
2	Smith	28000
3	Parker	35000
4	Smith	24000

Works_in	
employee_id	department_id
1	1
2	1
2	2
3	3
4	3

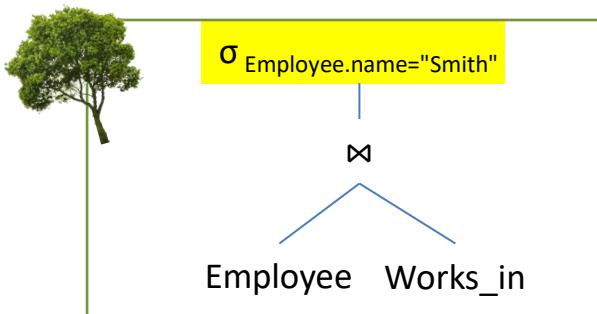
Equivalence rules

$$r \bowtie s = \pi_{R \cup S}(\sigma_{r.A_1 = s.A_1 \wedge r.A_2 = s.A_2 \wedge \dots \wedge r.A_n = s.A_n}(r \times s))$$

- Rule 5. The selection operation distributes over the natural join operation under the following two conditions
 - Rule 5a. It distributes when all the attributes in selection condition **involve only the attributes of one of the expressions** (say, E_1) being joined.

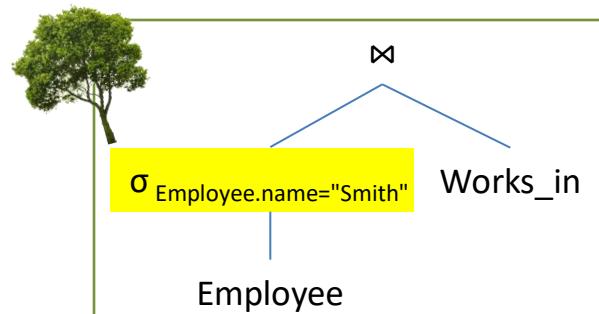
$$\sigma_p(E_1 \bowtie E_2) = (\sigma_p(E_1) \bowtie E_2)$$

$\sigma_{Employee.name='Smith'}($
Employee \bowtie Works_in
)



equivalent
to

$(\sigma_{Employee.name='Smith'}(Employee) \bowtie Works_in)$

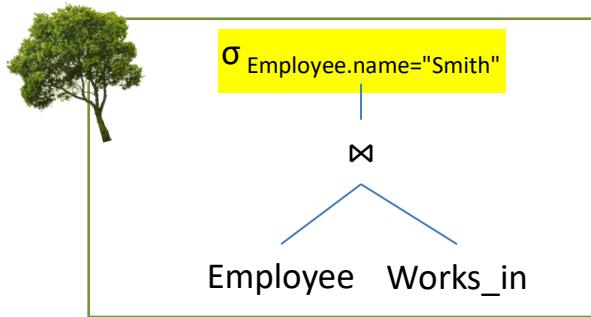


Equivalence rules



Expression tree A.

$\sigma_{Employee.name='Smith'}(Employee \bowtie Works_in)$



$\sigma_{Employee.name='Smith'}(Employee \bowtie Works_in)$

employee_id	name	salary	department_id
2	Smith	28000	1
2	Smith	28000	2
4	Smith	24000	3



Employee \bowtie Works_in

employee_id	name	salary	department_id
1	Jones	26000	1
2	Smith	28000	1
2	Smith	28000	2
3	Parker	35000	3
4	Smith	24000	3



Natural join evaluates $4 * 5 = 20$ combinations
result temporary relation consists of 5 tuples and 4 columns.

Employee

employee_id	name	salary
1	Jones	26000
2	Smith	28000
3	Parker	35000
4	Smith	24000

Works_in

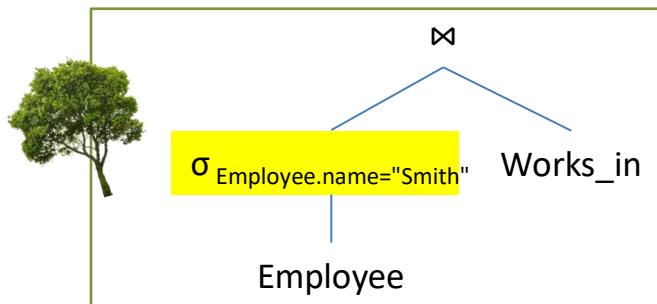
employee_id	department_id
1	1
2	1
2	2
3	3
4	3

Equivalence rules



Expression tree B.

$\sigma_{Employee.name='Smith'} (Employee) \bowtie Works_in$



$\sigma_{Employee.name='Smith'} (Employee) \bowtie Works_in$

employee_id	name	salary	department_id
2	Smith	28000	1
2	Smith	28000	2
4	Smith	24000	3



Natural join evaluates $2 * 5 = 10$ combinations
result temporary relation consists of 3 tuples and 4 columns.

$\sigma_{Employee.name='Smith'} (Employee)$

employee_id	name	salary
2	Smith	28000
4	Smith	24000



When comparing with the equivalence expression tree 1, we can see that if we push the selection predicates down the natural join (perform selection earlier than joining), then we will probably have a **smaller temporary relation**.

Employee

employee_id	name	salary
1	Jones	26000
2	Smith	28000
3	Parker	35000
4	Smith	24000

Works_in

employee_id	department_id
1	1
2	1
2	2
3	3
4	3

Equivalence rules

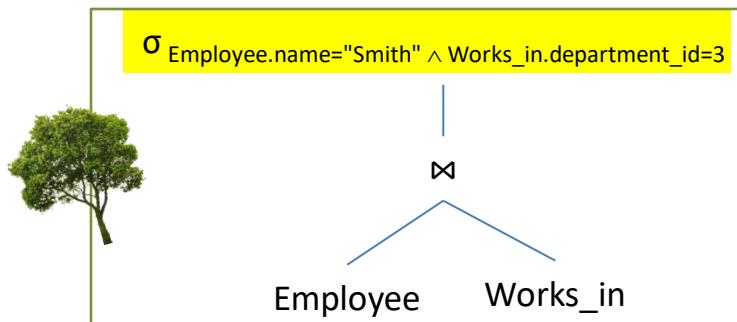
$$r \bowtie s = \pi_{R \cup S}(\sigma_{r.A1 = s.A1 \wedge r.A2 = s.A2 \wedge \dots \wedge r.An = s.An}(r \times s))$$

- Rule 5b. The selection distributes when selection condition p1 involves only the attributes of E_1 and p2 involves only the attributes of E_2 .

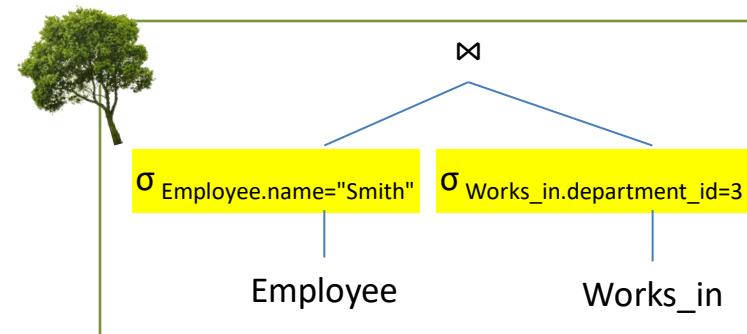
$$\sigma_{p1 \wedge p2}(E_1 \bowtie E_2) = (\sigma_{p1}(E_1) \bowtie \sigma_{p2}(E_2))$$

$\sigma_{\text{Employee.name}=\text{"Smith"} \wedge \text{Works_in.department_id}=3}$
Employee \bowtie **Works_in**
)

($\sigma_{\text{Employee.name}=\text{"Smith"}}(\text{Employee})$
 \bowtie
 $\sigma_{\text{Works_in.department_id}=3}(\text{Works_in})$)



equivalent
to

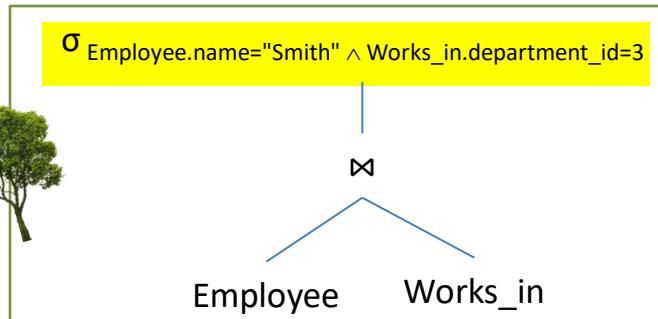


Equivalence rules



Expression tree A.

$\sigma_{Employee.name='Smith' \wedge Works_in.department_id=3}$
Employee \bowtie Works_in
)



$\sigma_{Employee.name='Smith' \wedge Works_in.department_id=3}$ (Employee \bowtie Works_in)

employee_id	name	salary	department_id
4	Smith	24000	3



Employee \bowtie Works_in

employee_id	name	salary	department_id
1	Jones	26000	1
2	Smith	28000	1
2	Smith	28000	2
3	Parker	35000	3
4	Smith	24000	3



Natural join evaluates $4 * 5 = 20$ combinations

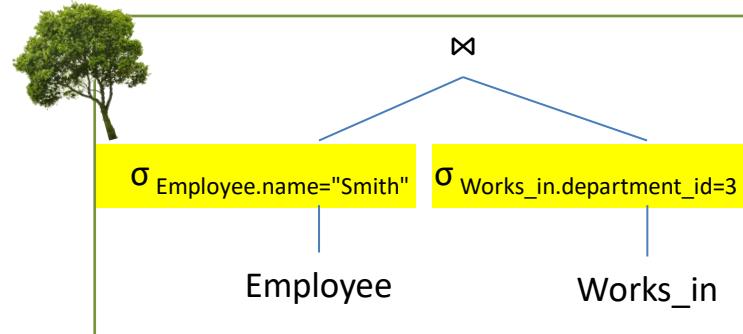
Employee			Works_in	
employee_id	name	salary	employee_id	department_id
1	Jones	26000	1	1
2	Smith	28000	2	1
3	Parker	35000	2	2
4	Smith	24000	3	3
			4	3

Equivalence rules



Expression tree B.

($\sigma_{Employee.name='Smith'}$ (Employee) \bowtie $\sigma_{Works_in.department_id=3}$ (Works_in))



($\sigma_{Employee.name='Smith'}$ (Employee) \bowtie $\sigma_{Works_in.department_id=3}$ (Works_in))

employee_id	name	salary	department_id
4	Smith	24000	3

Natural join evaluates 4 combinations

$\sigma_{Employee.name='Smith'}$ (Employee)

employee_id	name	salary
2	Smith	28000
4	Smith	24000

$\sigma_{Works_in.department_id=3}$ (Works_in)

employee_id	department_id
3	3
4	3

When comparing with the equivalence expression 1, we can see that if we push the selection predicates down the natural join (perform selection earlier than joining), the **natural join would consider fewer combinations**.



Employee			Works_in	
employee_id	name	salary	employee_id	department_id
1	Jones	26000	1	1
2	Smith	28000	2	1
3	Parker	35000	2	2
4	Smith	24000	3	3
			4	3

Equivalence rules

$$r \bowtie s = \pi_{R \cup S}(\sigma_{r.A_1 = s.A_1 \wedge r.A_2 = s.A_2 \wedge \dots \wedge r.A_n = s.A_n}(r \times s))$$

- Rule 6. The projection operation can distribute over the natural join operation.

$$\pi_{L_1 \cup L_2}(E_1 \bowtie E_2) = \pi_{L_1 \cup L_2}((\pi_{L_1 \cup L_3}(E_1)) \bowtie (\pi_{L_2 \cup L_3}(E_2)))$$

- Let L1 and L2 be some attributes from E1 and E2, respectively.
- Let L3 be attributes that are involved in join condition, but are not in L1 \cup L2.

Equivalence rules

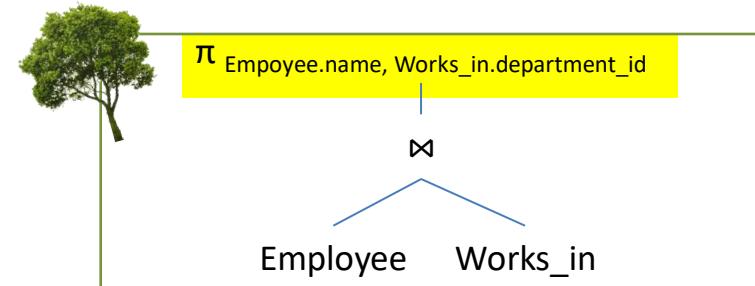
$$\pi_{L1 \cup L2} (E_1 \bowtie E_2) = \pi_{L1 \cup L2} ((\pi_{L1 \cup L3} (E_1)) \bowtie (\pi_{L2 \cup L3} (E_2)))$$

$\pi_{Employee.name, Works_in.department_id}$ (**Employee** \bowtie **Works_in**)

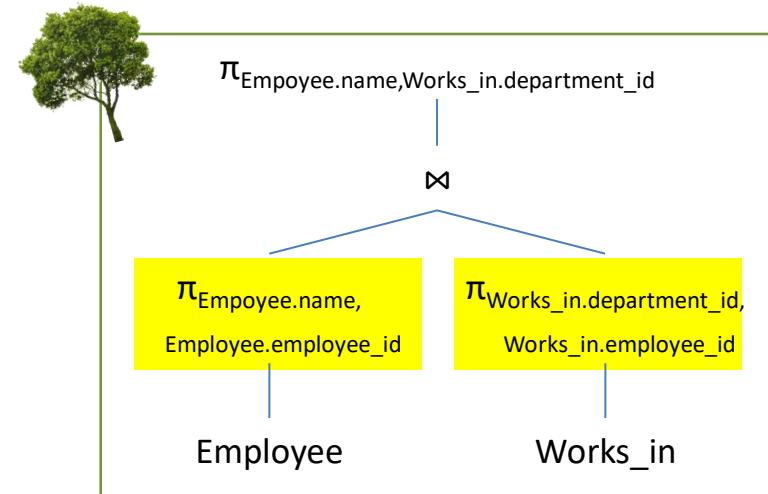
- L1 = Employee.name
- L2 = Works_in.department_id
- L3 = employee_id

The attribute used in natural join

$\pi_{Employee.name, Works_in.department_id} ($
 $\pi_{Employee.name, Employee.employee_id}$ (**Employee**)
 \bowtie
 $\pi_{Works_in.department_id, Works_in.employee_id}$ (**Works_in**)
)



equivalent to



Equivalence rules



Expression tree A.

$\pi_{Employee.name, Works_in.department_id}$
(Employee \bowtie Works_in)



$\pi_{Employee.name, Works_in.department_id}$

\bowtie

Employee Works_in

$\pi_{Employee.name, Works_in.department_id}$ **(Employee \bowtie Works_in)**

name	department_id
Jones	1
Smith	1
Smith	2
Parker	3
Smith	3

Employee \bowtie Works_in



employee_id	name	salary	department_id	since
1	Jones	26000	1	2012/1/1
2	Smith	28000	1	2011/3/2
2	Smith	28000	2	2014/2/1
3	Parker	35000	3	2013/2/2
4	Smith	24000	3	2013/2/8



Natural join evaluates $4 \times 5 = 20$ combinations,
result temporary relation consists of 5 columns.

Employee

Works_in

employee_id	name	salary
1	Jones	26000
2	Smith	28000
3	Parker	35000
4	Smith	24000

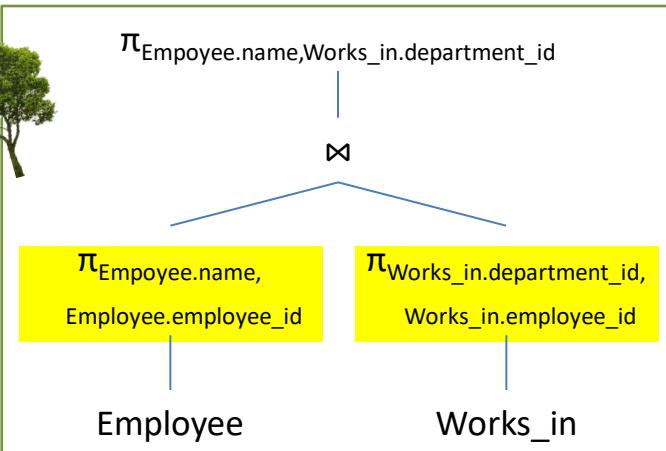
employee_id	department_id	since
1	1	2012/1/1
2	1	2011/3/2
2	2	2014/2/1
3	3	2013/2/2
4	3	2013/2/8

Equivalence rules



Expression tree B.

$\pi_{Employee.name, Works_in.department_id} ($
 $\pi_{Employee.name, Employee.employee_id} (Employee)$
 \bowtie
 $\pi_{Works_in.department_id, Works_in.employee_id} (Works_in)$
 $)$



$\pi_{Employee.name, Works_in.department_id} (\pi_{Employee.name, Employee.employee_id} (Employee) \bowtie \pi_{Works_in.department_id, Works_in.employee_id} (Works_in))$

name	department_id
Jones	1
Smith	1
Smith	2
Parker	3
Smith	3

$\pi_{Employee.name, Employee.employee_id} (Employee) \bowtie \pi_{Works_in.department_id, Works_in.employee_id} (Works_in)$

employee_id	name	department_id
1	Jones	1
2	Smith	1
2	Smith	2
3	Parker	3
4	Smith	3

Natural join evaluates $4 * 5 = 20$ combinations,
result temporary relation consists of 3 columns.

$\pi_{Employee.name, Employee.employee_id} (Employee) \bowtie \pi_{Works_in.department_id, Works_in.employee_id} (Works_in)$

employee_id	name	employee_id	name	employee_id	department_id
1	Jones	1	Jones	1	1
2	Smith	2	Smith	2	1
3	Parker	3	Parker	2	2
4	Smith	4	Smith	3	3
				4	3

Employee

employee_id	name	salary
1	Jones	26000
2	Smith	28000
3	Parker	35000
4	Smith	24000

Works_in

employee_id	department_id	since
1	1	2012/1/1
2	1	2011/3/2
2	2	2014/2/1
3	3	2013/2/2
4	3	2013/2/8

Equivalence rules

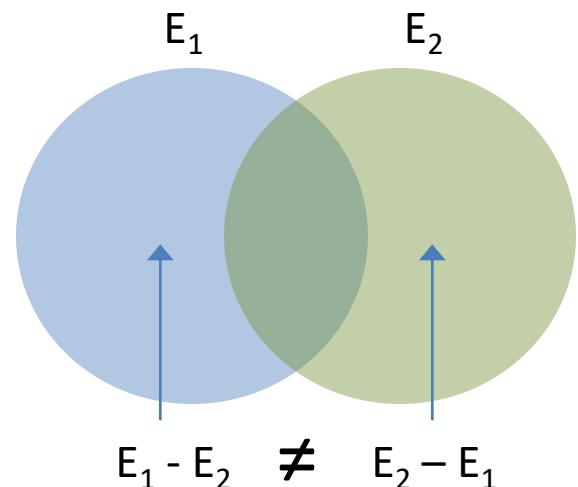
- Rule 7. The set operations union and intersections are commutative.

$$E_1 \cup E_2 = E_2 \cup E_1$$

$$E_1 \cap E_2 = E_2 \cap E_1$$

- The set different operation is NOT commutative

$$E_1 - E_2 \neq E_2 - E_1$$



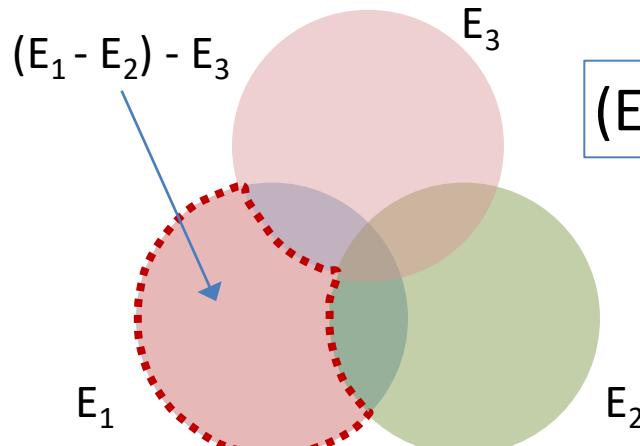
Equivalence rules

- Rule 8. The set operations union and intersections are associative.

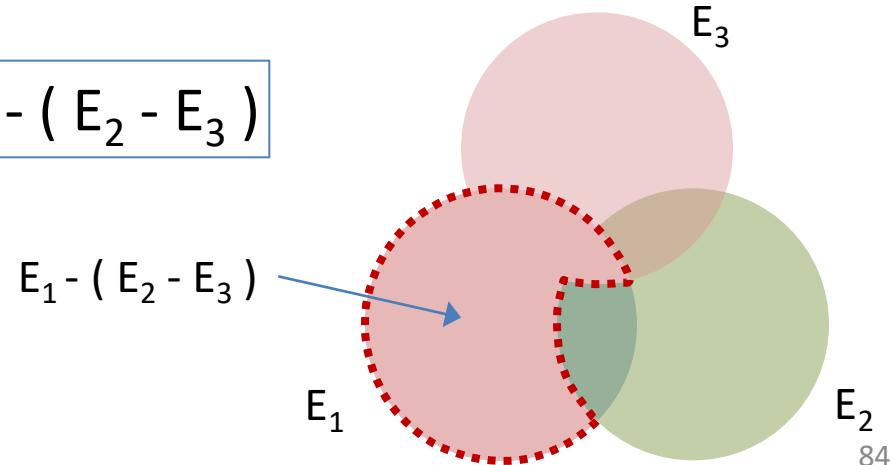
$$(E_1 \cup E_2) \cup E_3 = E_1 \cup (E_2 \cup E_3)$$

$$(E_1 \cap E_2) \cap E_3 = E_1 \cap (E_2 \cap E_3)$$

- The set different operation is NOT associative.



$$(E_1 - E_2) - E_3 \neq E_1 - (E_2 - E_3)$$



Equivalence rules



Rule 9. The selection operation distributes over the union, intersection and set difference operations

$$\sigma_p(E_1 \cup E_2) = \sigma_p(E_1) \cup \sigma_p(E_2)$$

$$\sigma_p(E_1 \cap E_2) = \sigma_p(E_1) \cap \sigma_p(E_2)$$

$$\sigma_p(E_1 - E_2) = \sigma_p(E_1) - \sigma_p(E_2)$$

Audio_CD

ID	name	provider_id	stock	#tracks
CD1	One Heart	P1	55	14
CD2	Miracle	P2	4	14

DVD

ID	name	provider_id	stock	length
DVD1	Prince of Persia	P2	3	110
DVD2	Iron man 3	P3	60	90
DVD3	Legend is born: Ip Man	P3	17	90

$$\sigma_{\text{stock} < 10} ($$

$$\pi_{\text{name}, \text{provider_id}, \text{stock}} (\text{Audio_CD}) \cup$$

$$\pi_{\text{name}, \text{provider_id}, \text{stock}} (\text{DVD})$$

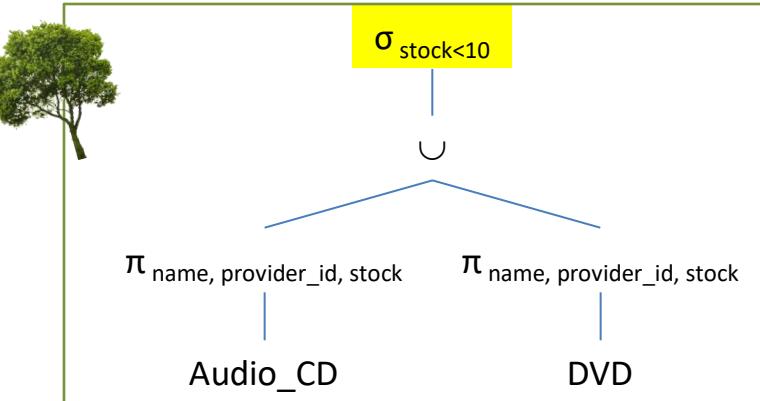
)

equivalent to

$$\sigma_{\text{stock} < 10} (\pi_{\text{name}, \text{provider_id}, \text{stock}} (\text{Audio_CD})) \cup$$

$$\sigma_{\text{stock} < 10} (\pi_{\text{name}, \text{provider_id}, \text{stock}} (\text{DVD}))$$

Equivalence rules

$$\sigma_{\text{stock} < 10} (\pi_{\text{name}, \text{provider_id}, \text{stock}} (\text{Audio_CD}) \cup \pi_{\text{name}, \text{provider_id}, \text{stock}} (\text{DVD}))$$


Audio_CD

ID	name	provider_id	stock	#tracks
CD1	One Heart	P1	55	14
CD2	Miracle	P2	4	14

DVD

ID	name	provider_id	stock	length
DVD1	Prince of Persia	P2	3	110
DVD2	Iron man 3	P3	60	90
DVD3	Legend is born: Ip Man	P3	17	90

$$\sigma_{\text{stock} < 10} (\pi_{\text{name}, \text{provider_id}, \text{stock}} (\text{Audio_CD}) \cup \pi_{\text{name}, \text{provider_id}, \text{stock}} (\text{DVD}))$$

name	provider_id	stock
Miracle	P2	4
Prince of Persia	P2	3

$$\pi_{\text{name}, \text{provider_id}, \text{stock}} (\text{Audio_CD}) \cup \pi_{\text{name}, \text{provider_id}, \text{stock}} (\text{DVD})$$

name	provider_id	stock
One Heart	P1	55
Miracle	P2	4
Prince of Persia	P2	3
Iron man 3	P3	60
Legend is born: Ip Man	P3	17

$$\pi_{\text{name}, \text{provider_id}, \text{stock}} (\text{Audio_CD})$$

name	provider_id	stock
One Heart	P1	55
Miracle	P2	4

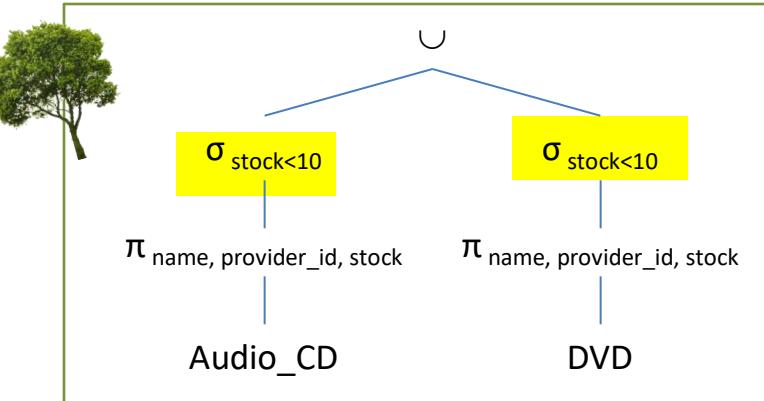
$$\pi_{\text{name}, \text{provider_id}, \text{stock}} (\text{DVD})$$

name	provider_id	stock
Prince of Persia	P2	3
Iron man 3	P3	60
Legend is born: Ip Man	P3	17

Equivalence rules

$\sigma_{\text{stock} < 10}(\pi_{\text{name}, \text{provider_id}, \text{stock}}(\text{Audio_CD})) \cup \sigma_{\text{stock} < 10}(\pi_{\text{name}, \text{provider_id}, \text{stock}}(\text{DVD}))$

$\sigma_{\text{stock} < 10}(\pi_{\text{name}, \text{provider_id}, \text{stock}}(\text{Audio_CD})) \cup \sigma_{\text{stock} < 10}(\pi_{\text{name}, \text{provider_id}, \text{stock}}(\text{DVD}))$



Audio_CD

ID	name	provider_id	stock	#tracks
CD1	One Heart	P1	55	14
CD2	Miracle	P2	4	14

DVD

ID	name	provider_id	stock	length
DVD1	Prince of Persia	P2	3	110
DVD2	Iron man 3	P3	60	90
DVD3	Legend is born: Ip Man	P3	17	90

name	provider_id	stock
Miracle	P2	4
Prince of Persia	P2	3

$\sigma_{\text{stock} < 10}(\pi_{\text{name}, \text{provider_id}, \text{stock}}(\text{Audio_CD}))$

name	provider_id	stock
Miracle	P2	4

$\sigma_{\text{stock} < 10}(\pi_{\text{name}, \text{provider_id}, \text{stock}}(\text{DVD}))$

name	provider_id	stock
Prince of Persia	P2	3

$\pi_{\text{name}, \text{provider_id}, \text{stock}}(\text{Audio_CD})$

name	provider_id	stock
One Heart	P1	55
Miracle	P2	4

$\pi_{\text{name}, \text{provider_id}, \text{stock}}(\text{DVD})$

name	provider_id	stock
Prince of Persia	P2	3
Iron man 3	P3	60
Legend is born: Ip Man	P3	17

Equivalence rules

- Rule 10. The projection operation distributes over the union operation

$$\pi_L(E_1 \cup E_2) = \pi_L(E_1) \cup \pi_L(E_2)$$

- Projection does not distribute over intersection and set difference.

$$\pi_L(E_1 \cap E_2) \neq \pi_L(E_1) \cap \pi_L(E_2)$$

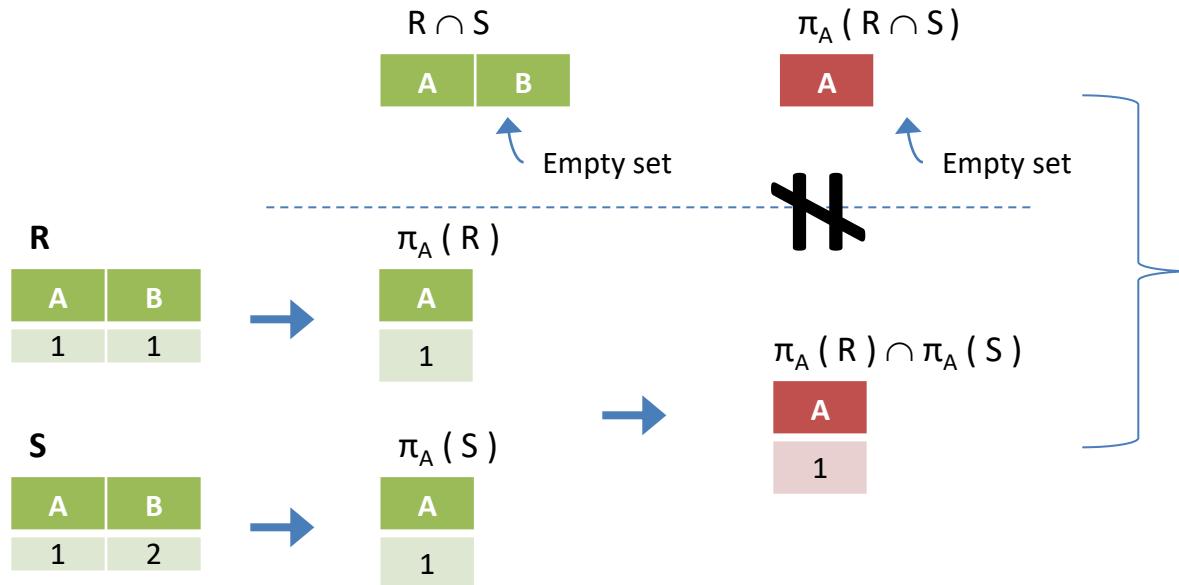
$$\pi_L(E_1 - E_2) \neq \pi_L(E_1) - \pi_L(E_2)$$

Equivalence rules



Why projection does not distribute over intersection and set difference?

To show that projection does not distribute over intersection, you only need to provide a counter example.



This counter example shows that
 $\pi_A(R \cap S) \neq \pi_A(R) \cap \pi_A(S)$

Now, can you try to construct a counter example to show that projection does not distribute over set difference?

Section 6

Example of
query optimization

Transformation

- Find the names of all instructors in the CS department ($dpt_id = 1$) who have taught a course in 2nd semester, together with the course title of all the courses that the instructors teach.

```
SELECT I.name, C.title  
FROM Instructor I, Teaches T, Course C  
WHERE I.dpt_id = 1 AND  
T.sem=2 AND  
I.instructor_id = T.instructor_id AND  
T.course_id = C.course_id ;
```

SQL

$$\pi_{I.name, C.title} (\sigma_{I.dpt_id = 1 \wedge T.sem=2} (\rho_I(\text{Instructor}) \bowtie (\rho_T(\text{Teaches}) \bowtie \rho_C(\text{Course}))))$$

Relational algebra

Instructor		
instructor_id	name	dpt_id
1	Kit	1
2	Ben	1
3	Michael	2
4	William	3

Teaches		
instructor_id	course_id	sem
1	1	1
1	2	2
2	4	1
3	3	2

Course		
course_id	title	credit
1	Intro to DB	6
2	Programming I	6
3	Accounting	6
4	Algorithms	6

Transformation

$$\pi_{I.name,C.title} \left(\sigma_{I.dpt_id=1 \wedge T.sem=2} \left(\rho_I(\text{Instructor}) \bowtie (\rho_T(\text{Teaches}) \bowtie \rho_C(\text{Course})) \right) \right)$$

equivalent to

$$\pi_{I.name,C.title} \left(\sigma_{I.dpt_id=1 \wedge T.sem=2} \left(\rho_I(\text{Instructor}) \bowtie \left(\rho_T(\text{Teaches}) \bowtie \pi_{C.course_id,C.title}(\rho_C(\text{Course})) \right) \right) \right)$$

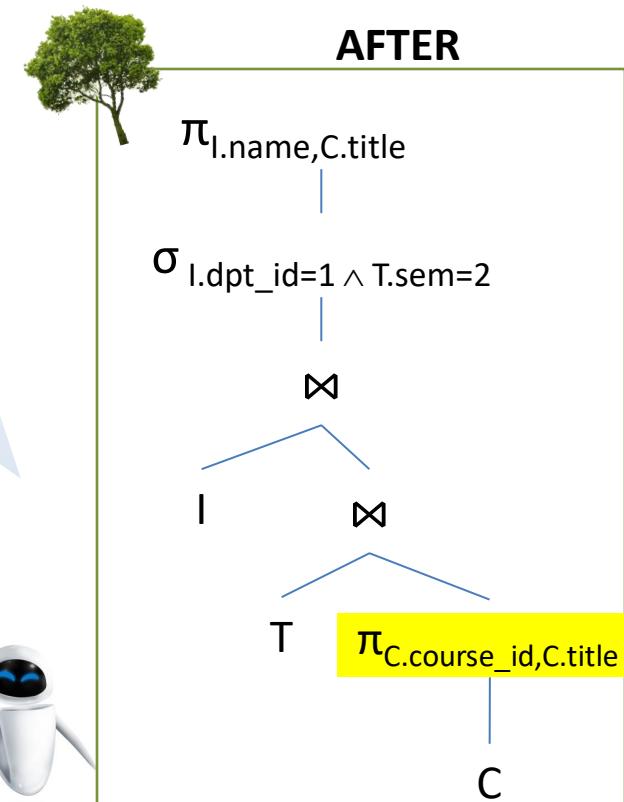
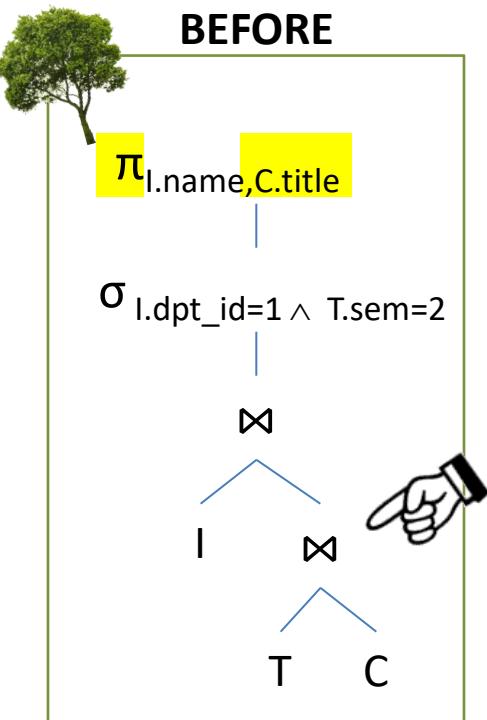
Rule 6

$$\pi_{L_1 \cup L_2}(E_1 \bowtie E_2) = \pi_{L_1 \cup L_2}((\pi_{L_1 \cup L_3}(E_1)) \bowtie (\pi_{L_2 \cup L_3}(E_2)))$$

Let's try to push the projection $\pi_{C.title}$ downward and apply it ahead of the natural joins.

Since this natural join requires $C.course_id = T.course_id$, therefore, **we have to add the joining attribute $C.course_id$ to make the projection**

$\pi_{C.course_id,C.title}$



Transformation

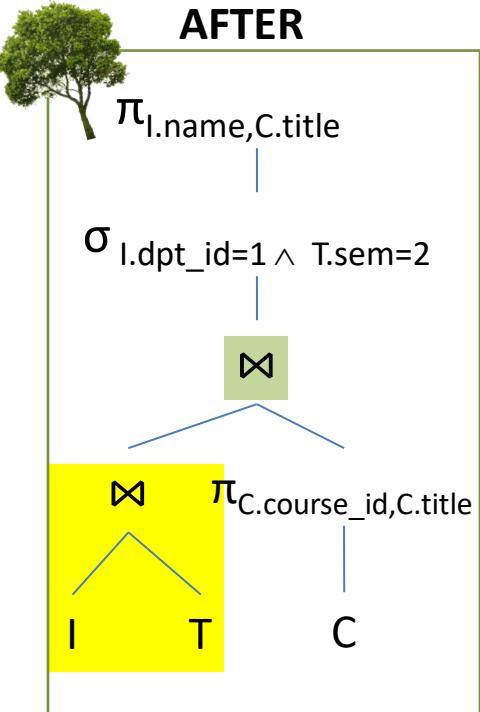
$$\pi_{I.name,C.title} \left(\sigma_{I.dpt_id=1 \wedge T.sem=2} \left(\rho_I(\text{Instructor}) \bowtie \rho_T(\text{Teaches}) \right) \bowtie \pi_{C.course_id, C.title} (\rho_C(\text{Course})) \right)$$


equivalent
to

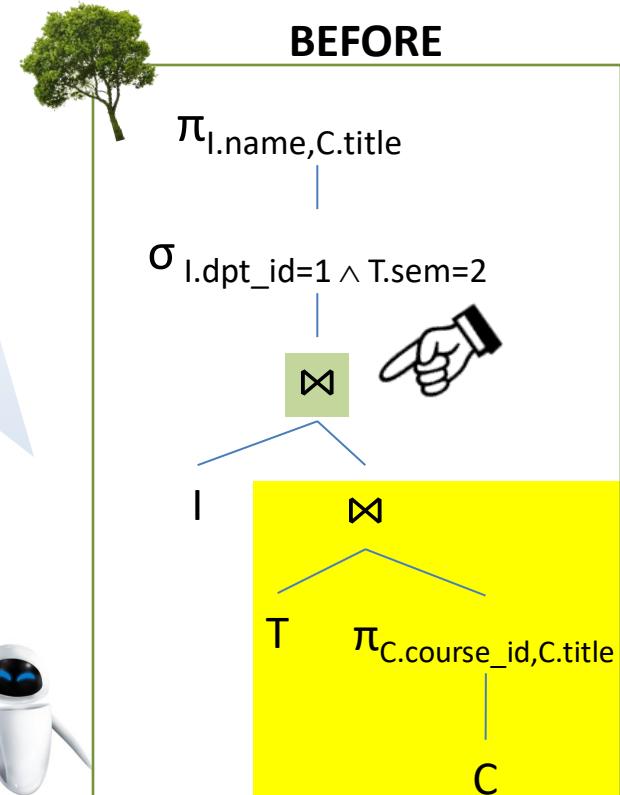
$$\pi_{I.name,C.title} \left(\sigma_{I.dpt_id=1 \wedge T.sem=2} \left(\rho_I(\text{Instructor}) \bowtie \left(\rho_T(\text{Teaches}) \bowtie \pi_{C.course_id, C.title} (\rho_C(\text{Course})) \right) \right) \right)$$

Rule 4

AFTER



BEFORE



- ($E_1 \bowtie E_2$) $\bowtie E_3 = E_1 \bowtie (E_2 \bowtie E_3)$
- Now we would like to push the selection down to reduce the size of the temporary result of the natural join.
 - As the selection involves relations I and T only, we would like to rearrange the natural joins to make I and T under one natural join.
 - Since natural joins are associative, we can make such rearrangement.

Transformation

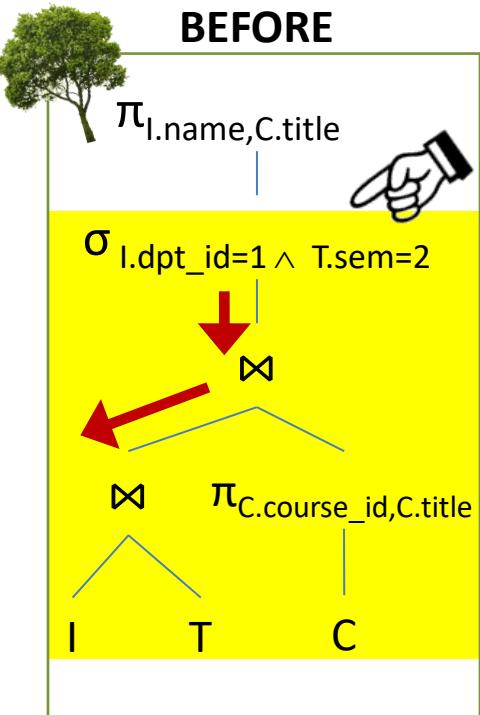
$$\begin{aligned} & \pi_{I.name, C.title} (\\ & \sigma_{I.dpt_id = 1 \wedge T.sem = 2} (\\ & \quad (\rho_I(\text{Instructor}) \bowtie \rho_T(\text{Teaches})) \\ & \quad \bowtie \pi_{C.course_id, C.title} (\rho_C(\text{Course})) \\ &) \\ &) \end{aligned}$$

equivalent to

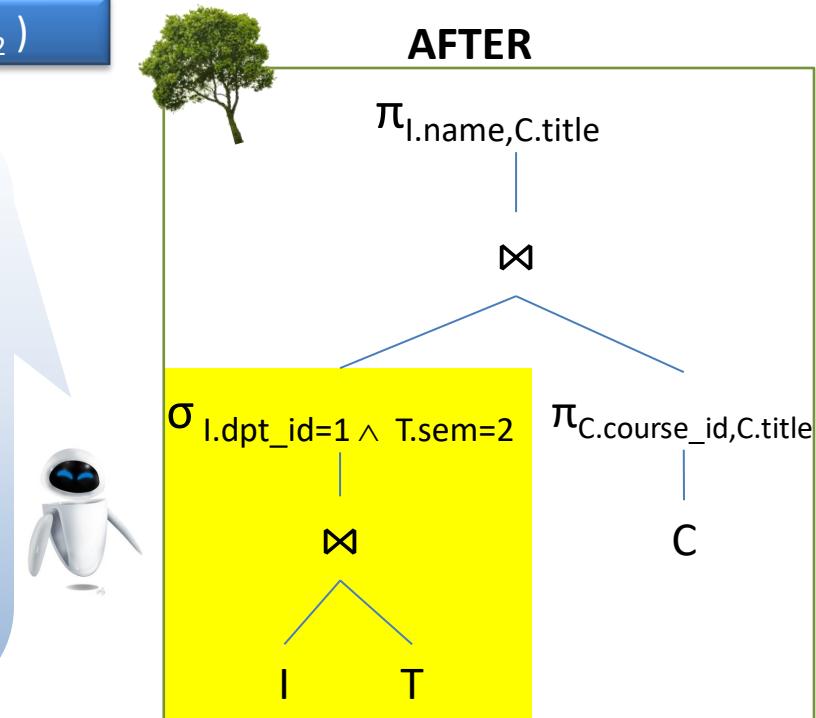
$$\begin{aligned} & \pi_{I.name, C.title} (\\ & \sigma_{I.dpt_id = 1 \wedge T.sem = 2} (\\ & \quad \rho_I(\text{Instructor}) \bowtie \rho_T(\text{Teaches}) \\ &) \\ &) \\ & \bowtie \pi_{C.course_id, C.title} (\rho_C(\text{Course})) \\ &) \end{aligned}$$

Rule 5a

$$\sigma_{p1}(E_1 \bowtie E_2) = (\sigma_{p1}(E_1) \bowtie E_2)$$



- ➊ Now we can push the selection down one level.
- ➋ According to Rule 5a, we can distribute both selection predicates to the L.H.S. of the selection as the R.H.S. does not contain any attribute in the selection predicate.



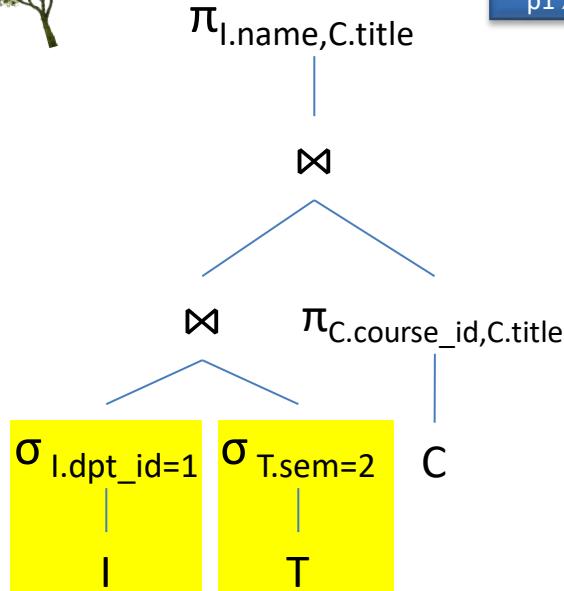
Transformation

$$\begin{aligned} \pi_{I.name, C.title} (& \\ & (\sigma_{I.dpt_id=1} (\rho_I(\text{Instructor})) \\ & \bowtie \\ & \sigma_{T.sem=2} (\rho_T(\text{Teaches})) \\) \\ & \bowtie \pi_{C.course_id, C.title} (\rho_C(\text{Course})) \\) \end{aligned}$$

equivalent
to

$$\begin{aligned} \pi_{I.name, C.title} (& \\ & (\sigma_{I.dpt_id=1 \wedge T.sem=2} (\rho_I(\text{Instructor}) \bowtie \rho_T(\text{Teaches})) \\) \\ & \bowtie \pi_{C.course_id, C.title} (\rho_C(\text{Course})) \\) \end{aligned}$$


AFTER



Rule 5b

$$\sigma_{p_1 \wedge p_2} (E_1 \bowtie E_2) = (\sigma_{p_1} (E_1) \bowtie \sigma_{p_1} (E_2))$$

Now we can further push the selection one more level down by applying rule 5b.

According to Rule 5b, we can distribute

- $\sigma_{I.dpt_id=1}$ to I
- $\sigma_{T.sem=2}$ to T



BEFORE

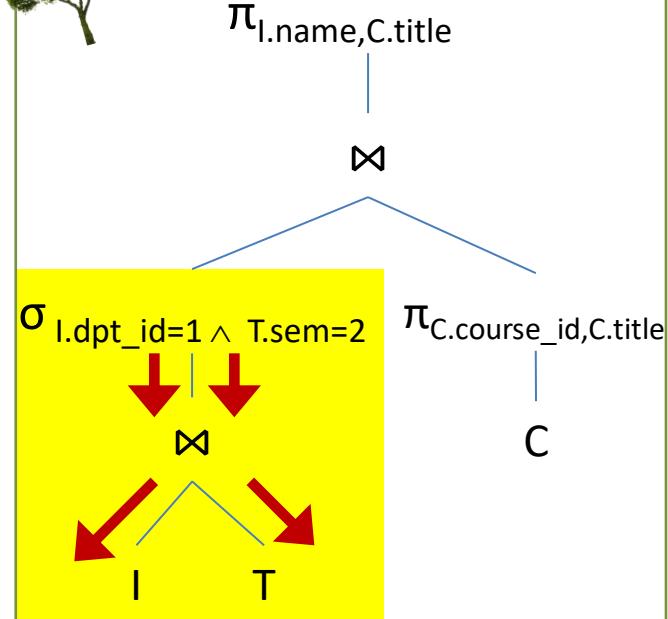
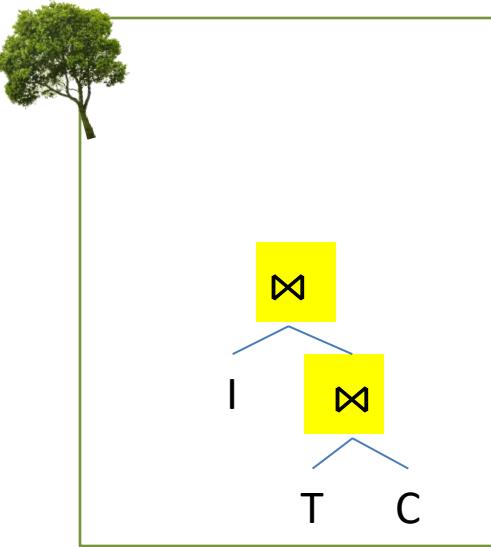


Illustration (original tree)

$$\pi_{I.name, C.title} \left(\sigma_{I.dpt_id = 1 \wedge T.sem = 2} \left(\rho_I(\text{Instructor}) \bowtie (\rho_T(\text{Teaches}) \bowtie \rho_C(\text{Course})) \right) \right)$$


Instructor

instructor_id	name	dpt_id
1	Kit	1
2	Ben	1
3	Michael	2
4	William	3

Teaches

instructor_id	course_id	sem
1	1	1
1	2	2
2	4	1
3	3	2

Course

course_id	title	credit
1	Intro to DB	6
2	Programming I	6
3	Accounting	6
4	Algorithms	6

$$\rho_I(\text{Instructor}) \bowtie (\rho_T(\text{Teaches}) \bowtie \rho_C(\text{Course}))$$

instructor_id	name	dpt_id	course_id	sem	title	credit
1	Kit	1	1	1	Intro to DB	6
1	Kit	1	2	2	Programming I	6
2	Ben	1	4	1	Algorithms	6
3	Michael	2	3	2	Accounting	6

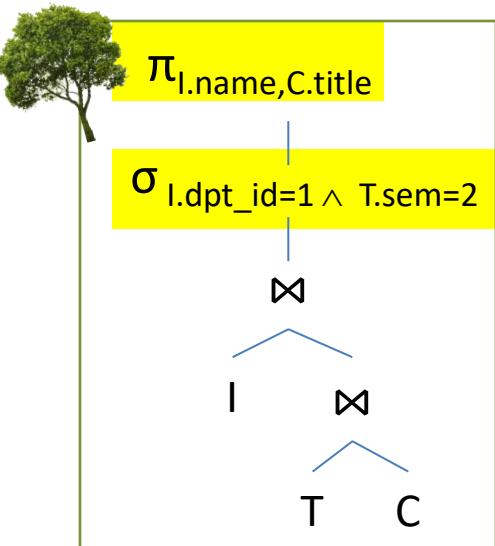
Natural join evaluates $4 * 4 = 16$ combinations.

$$\rho_T(\text{Teaches}) \bowtie \rho_C(\text{Course})$$

instructor_id	name	dpt_id	course_id	sem
1	Kit	1	1	1
1	Kit	1	2	2
2	Ben	1	4	1
3	Michael	2	3	2

Natural join evaluates $4 * 4 = 16$ combinations.

Illustration (original tree)

$$\pi_{I.name,C.title} \left(\sigma_{I.dpt_id=1 \wedge T.sem=2} \left(\rho_I(\text{Instructor}) \bowtie (\rho_T(\text{Teaches}) \bowtie \rho_C(\text{Course})) \right) \right)$$


$$\rho_I(\text{Instructor}) \bowtie (\rho_T(\text{Teaches}) \bowtie \rho_C(\text{Course}))$$

instructor_id	name	dpt_id	course_id	sem	title	credit
1	Kit	1	1	1	Intro to DB	6
1	Kit	1	2	2	Programming I	6
2	Ben	1	4	1	Algorithms	6
3	Michael	2	3	2	Accounting	6



$$\sigma_{I.dpt_id=1 \wedge T.sem=2} \left(\rho_I(\text{Instructor}) \bowtie (\rho_T(\text{Teaches}) \bowtie \rho_C(\text{Course})) \right)$$

instructor_id	name	dpt_id	course_id	sem	title	credit
1	Kit	1	2	2	Programming I	6



$$\pi_{I.name,C.title} \left(\sigma_{I.dpt_id=1 \wedge T.sem=2} \left(\rho_I(\text{Instructor}) \bowtie (\rho_T(\text{Teaches}) \bowtie \rho_C(\text{Course})) \right) \right)$$

name	title
Kit	Programming I

Illustration (transformed tree)

$$\begin{aligned} & \pi_{I.name, C.title} (\\ & \quad (\sigma_{I.dpt_id=1} (\rho_I(\text{Instructor})) \\ & \quad \bowtie \\ & \quad \sigma_{T.sem=2} (\rho_T(\text{Teaches})) \\ &) \\ & \bowtie \pi_{C.course_id, C.title} (\rho_C(\text{Course})) \end{aligned})$$

$$\sigma_{I.dpt_id=1} (\rho_I(\text{Instructor})) \bowtie \sigma_{T.sem=2} (\rho_T(\text{Teaches}))$$

instructor_id	name	dpt_id	course_id	sem
1	Kit	1	2	2

$$\sigma_{I.dpt_id=1} (\rho_I(\text{Instructor}))$$

instructor_id	name	dpt_id
1	Kit	1
2	Ben	1

Instructor

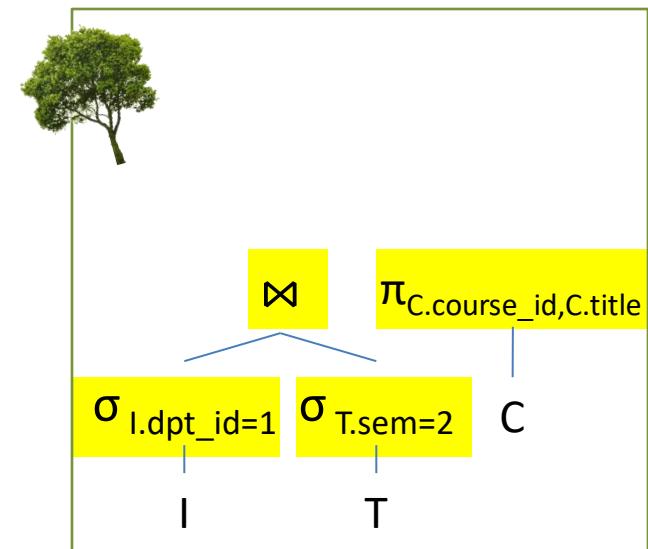
instructor_id	name	dpt_id
1	Kit	1
2	Ben	1
3	Michael	2
4	William	3

$$\sigma_{T.sem=2} (\rho_T(\text{Teaches}))$$

instructor_id	course_id	sem
1	1	1
1	2	2
2	4	1
3	3	2

Teaches

instructor_id	course_id	sem
1	1	1
1	2	2
2	4	1
3	3	2



$$\pi_{C.course_id, C.title} (\rho_C(\text{Course}))$$

course_id	title
1	Intro to DB
2	Programming I
3	Accounting
4	Algorithms

Course

course_id	title	credit
1	Intro to DB	6
2	Programming I	6
3	Accounting	6
4	Algorithms	6

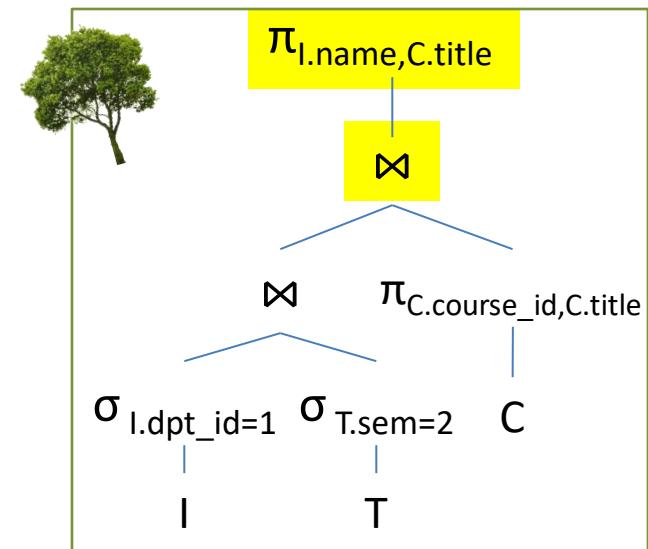
Illustration (transformed tree)

$$\begin{aligned} & \pi_{I.name, C.title} \left(\begin{array}{l} \sigma_{I.dpt_id=1} \left(\rho_I(\text{Instructor}) \right) \\ \bowtie \\ \sigma_{T.sem=2} \left(\rho_T(\text{Teaches}) \right) \end{array} \right) \\ & \bowtie \pi_{C.course_id, C.title} \left(\rho_C(\text{Course}) \right) \end{aligned}$$

$$\sigma_{I.dpt_id=1} \left(\rho_I(\text{Instructor}) \right) \bowtie \sigma_{T.sem=2} \left(\rho_T(\text{Teaches}) \right)$$

instructor_id	name	dpt_id	course_id	sem
1	Kit	1	2	2

Natural join evaluates $1 * 4 = 4$ combinations.



$$\pi_{C.course_id, C.title} \left(\rho_C(\text{Course}) \right)$$

course_id	title
1	Intro to DB
2	Programming I
3	Accounting
4	Algorithms

$$\sigma_{I.dpt_id=1} \left(\rho_I(\text{Instructor}) \right) \bowtie \sigma_{T.sem=2} \left(\rho_T(\text{Teaches}) \right) \bowtie \pi_{C.course_id, C.title} \left(\rho_C(\text{Course}) \right)$$

instructor_id	name	dpt_id	course_id	sem	title
1	Kit	1	2	2	Programming I

$$\pi_{I.name, C.title} \left(\sigma_{I.dpt_id=1} \left(\rho_I(\text{Instructor}) \right) \bowtie \sigma_{T.sem=2} \left(\rho_T(\text{Teaches}) \right) \bowtie \pi_{C.course_id, C.title} \left(\rho_C(\text{Course}) \right) \right)$$

name	title
Kit	Programming I

Summary

- **Relational algebra (RA) defines a set of algebraic operations on tables, and output tables as result.**
 - 6 fundamental operations
 - Additional operations does not extend the power of the fundamental operators, but they simplify the expression.
 - Extended operations add expressive power.
- **Relational algebra (RA) is the basics of query optimization.**

Lecture 6

END

COMP3278

Introduction to Database Management Systems

Dr. Ping Luo

Email : pluo@cs.hku.hk



Department of Computer Science, The University of Hong Kong